

- **Think** till **left fork** is available; when it is → **Pick it up**
- **Think** till **right fork** is available; when it is → **Pick it up**
- **Eat** for a fixed amount of time when **both forks** are held
- **Put the right fork down** → **Put the left fork down**
- **Repeat** the process


**DEADLOCK**

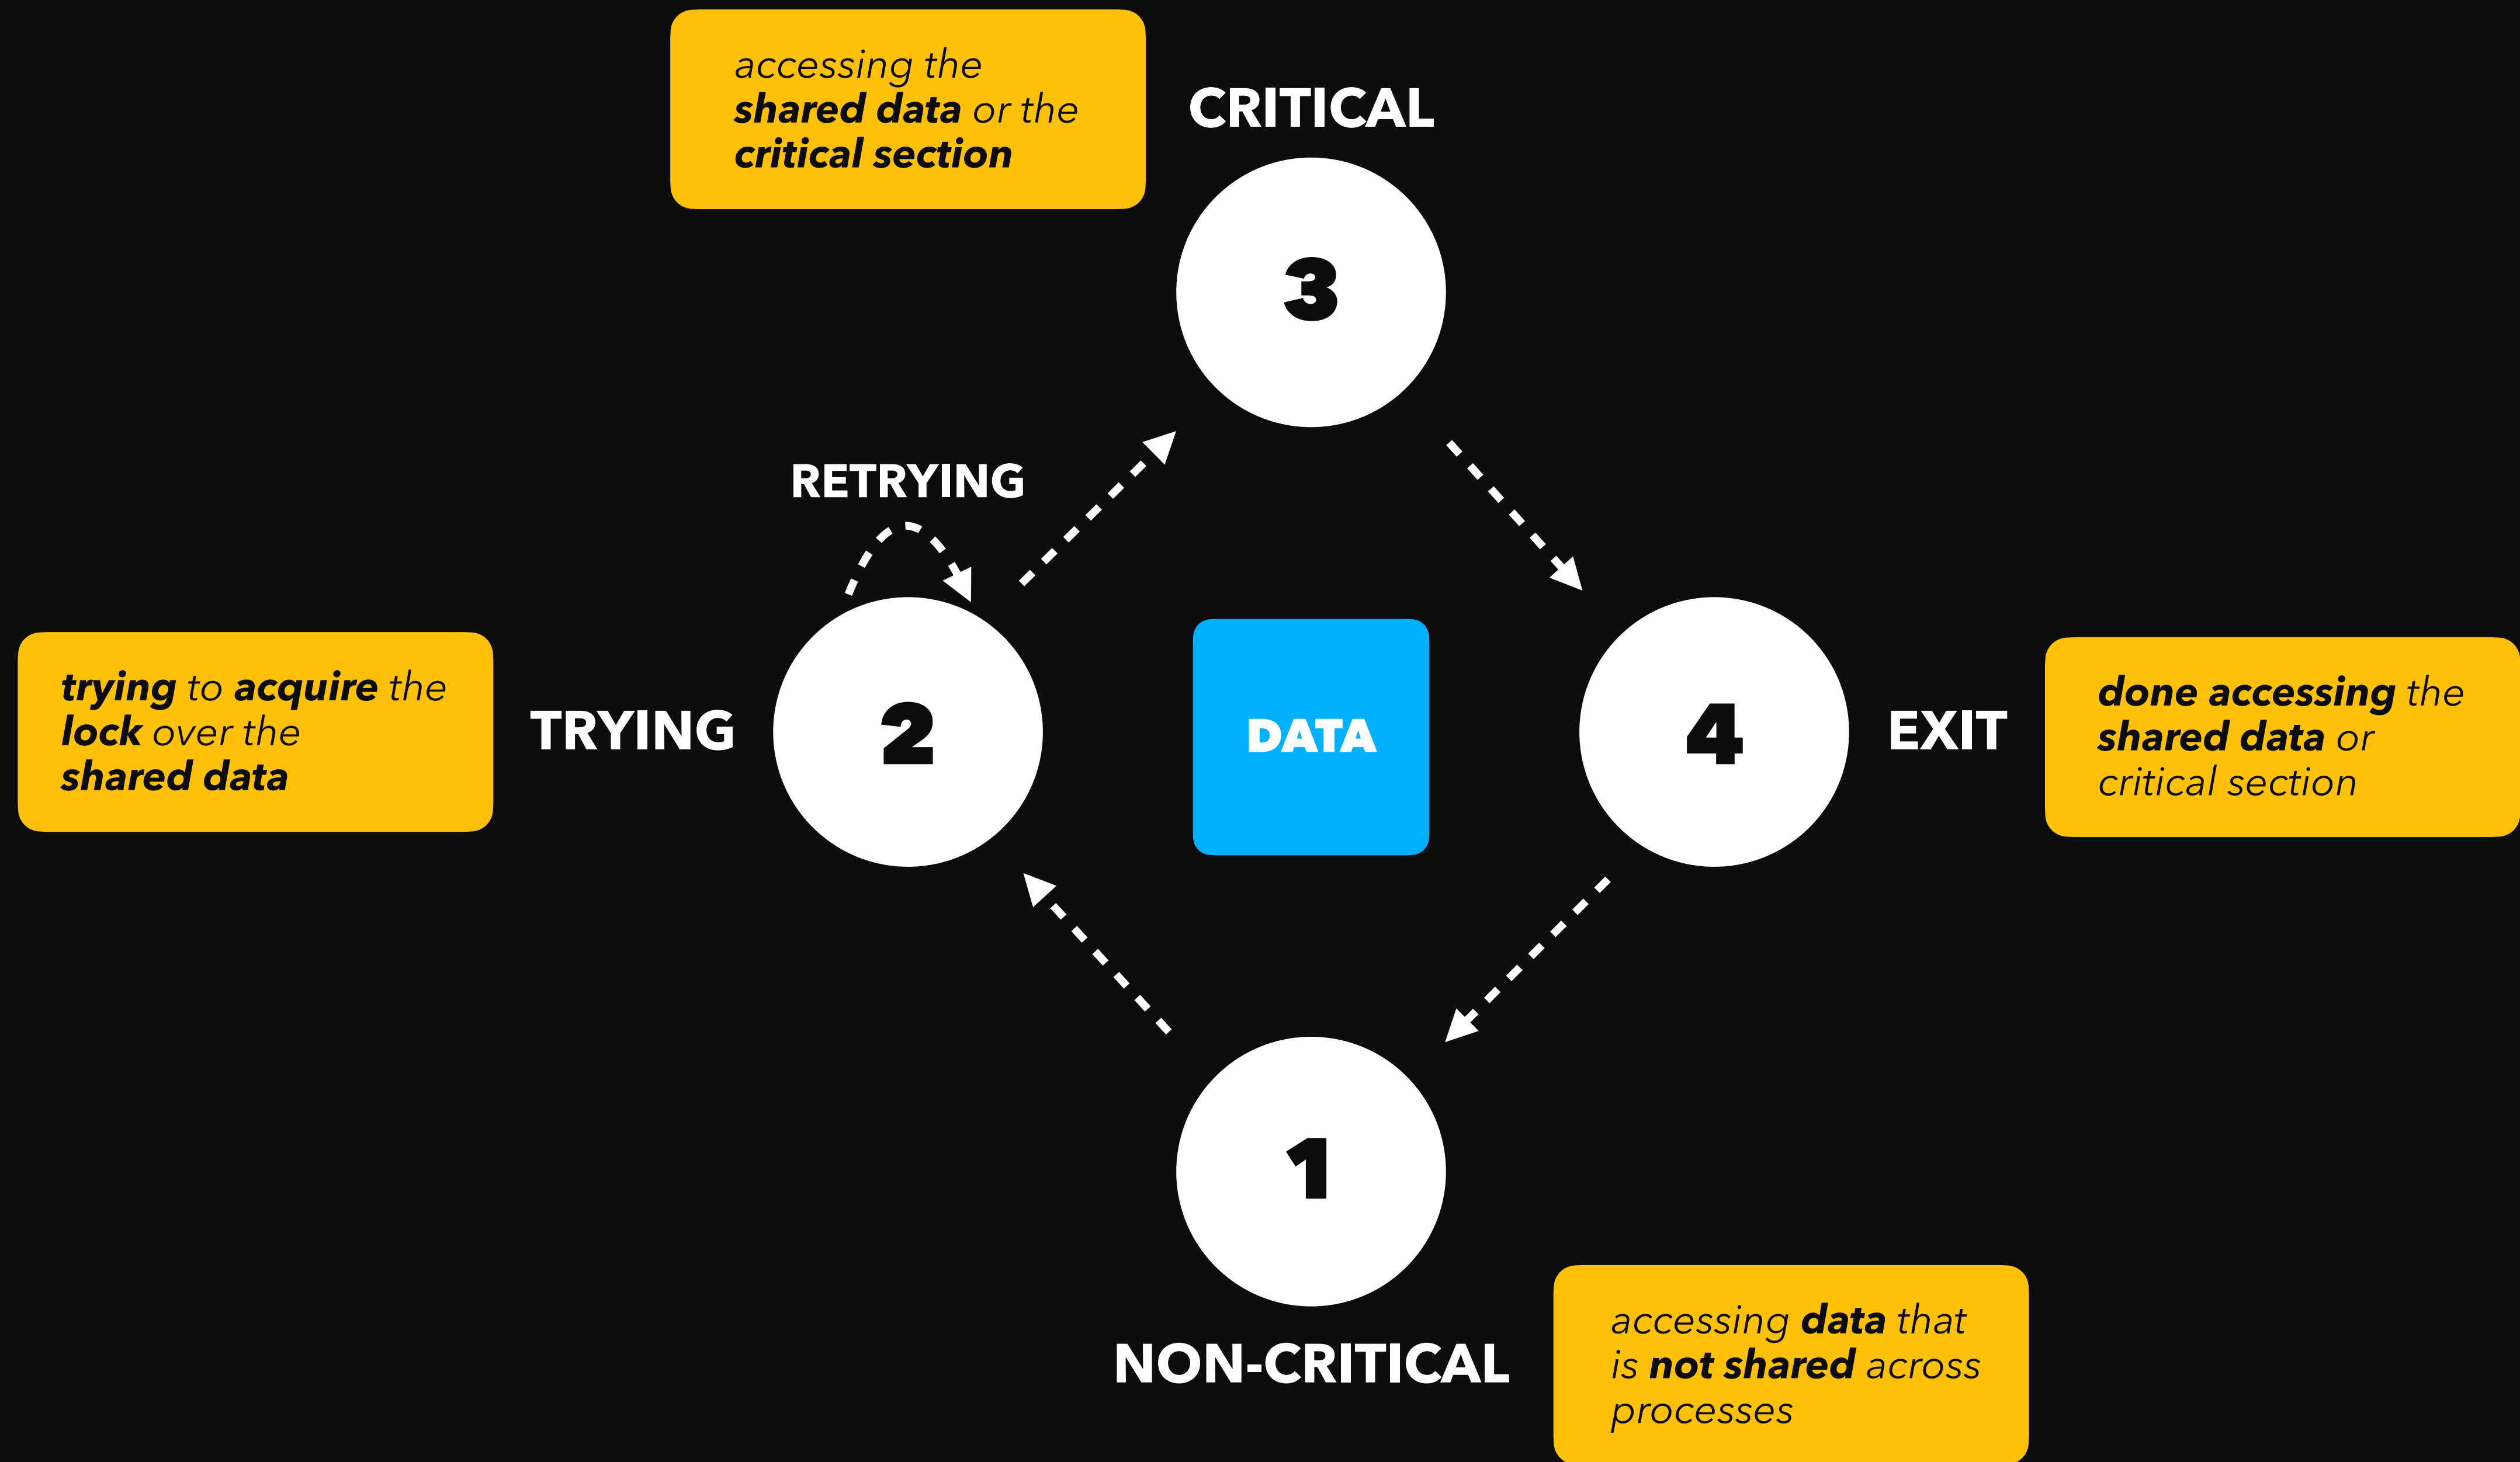
**STARVATION**

**LIVELOCK**

**MUTUAL EXCLUSION**

In computer science **Mutual Exclusion** is a property of **Concurrency Control**, which is instituted for the purpose of preventing **Race Conditions**.





LOCKS	MUTEX
READERS-WRITER LOCKS	RWMutex
RECURSIVE LOCKS	UNAVAILABLE
SEMAPHORES	INTERNAL
MONITORS	INTERNAL



1




2



3






**ORDER**

**RESULT**

**CORRECTNESS**





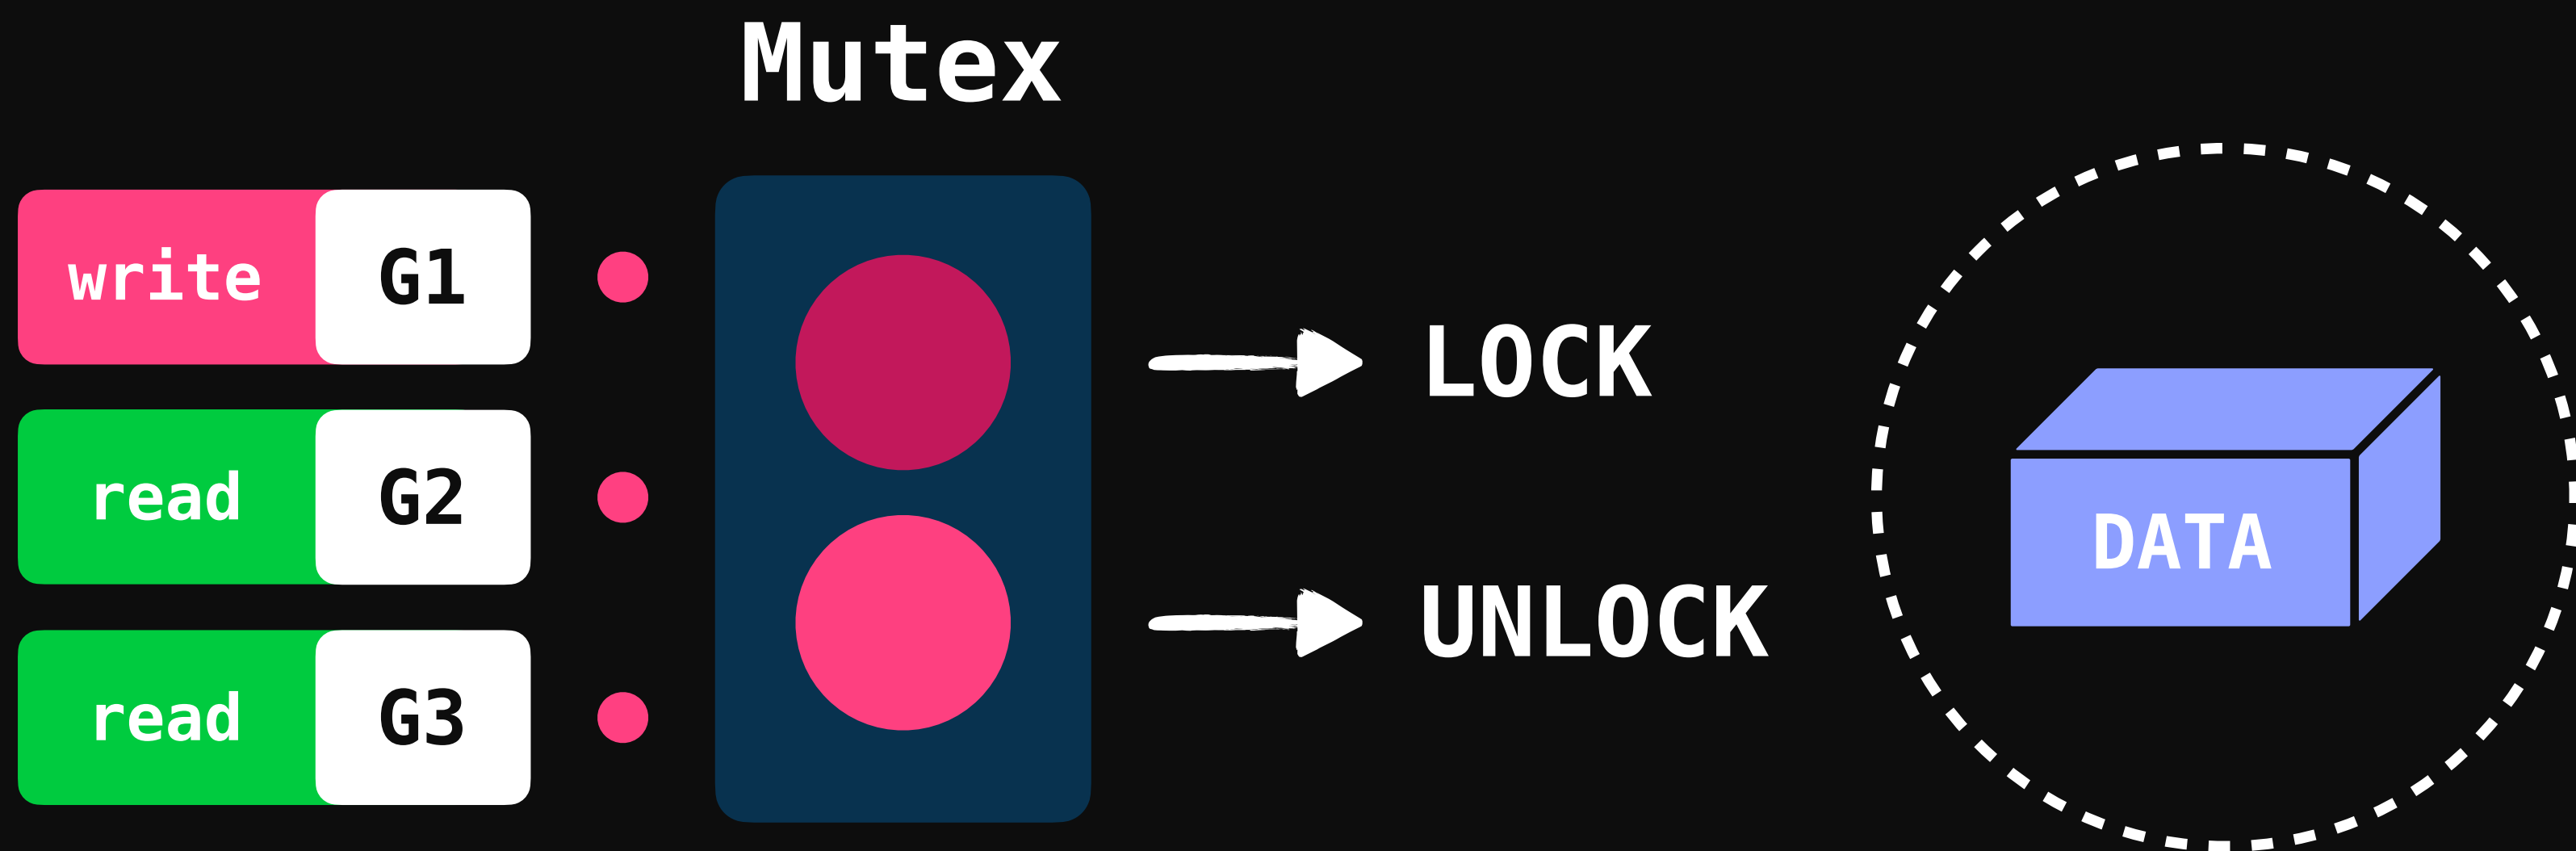


**READ WRITE ONCE**

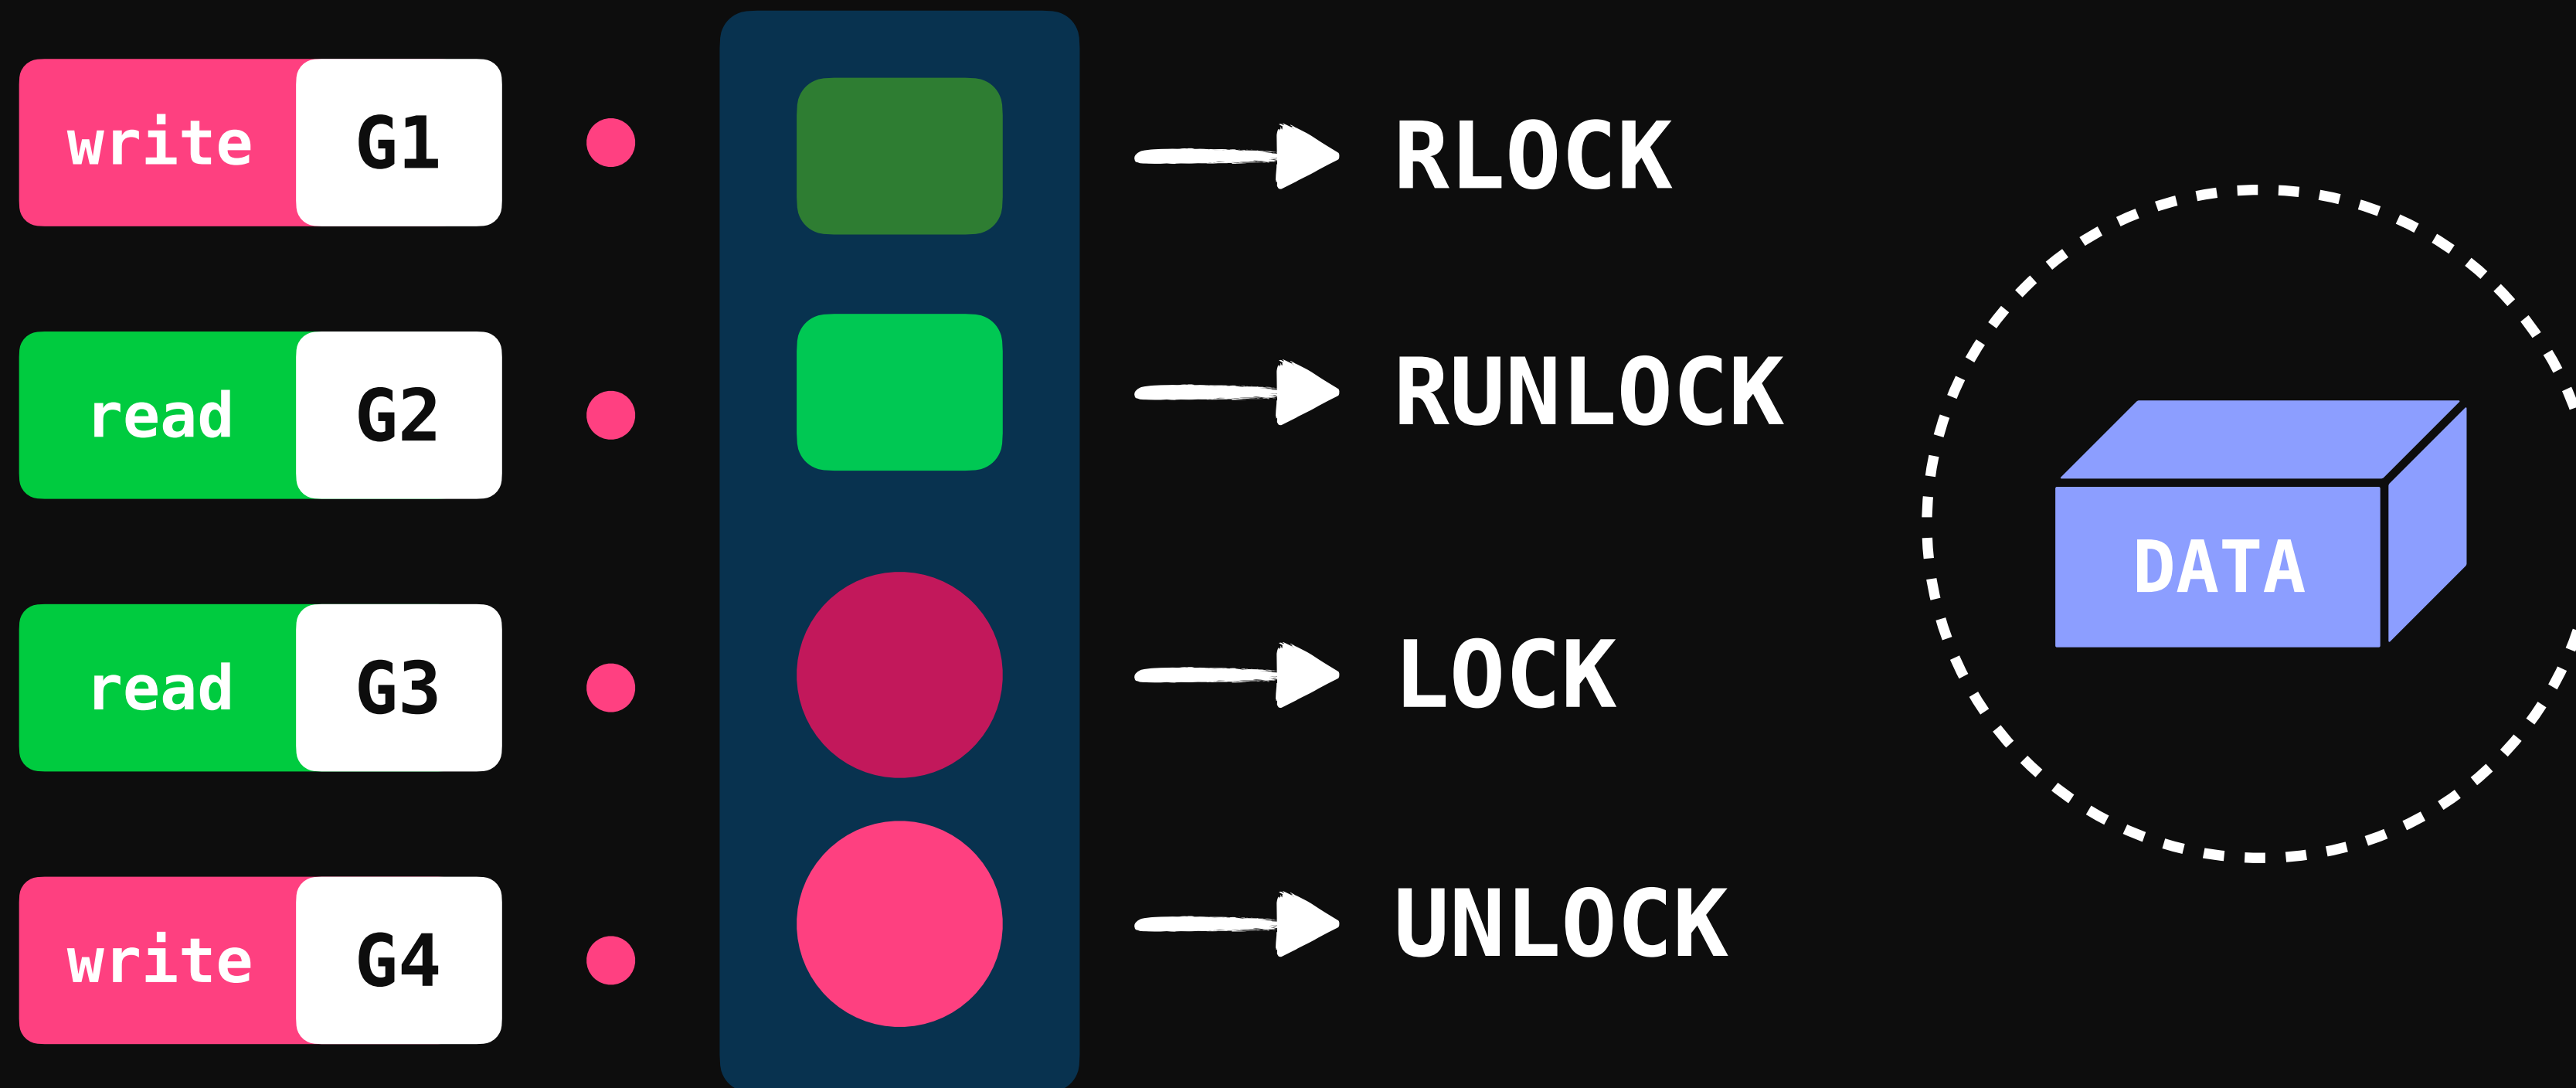
**MUTEX**

**WRITE ONCE, READ MANY**

**RWMUTEX**



## RWMutex



**G1**

**G2**

**G3**

`i++`

`i++`

`i++`

`var i`

`i` could be 1

`i` could be 2

`i` could be 3



`i++`

`get value of i`

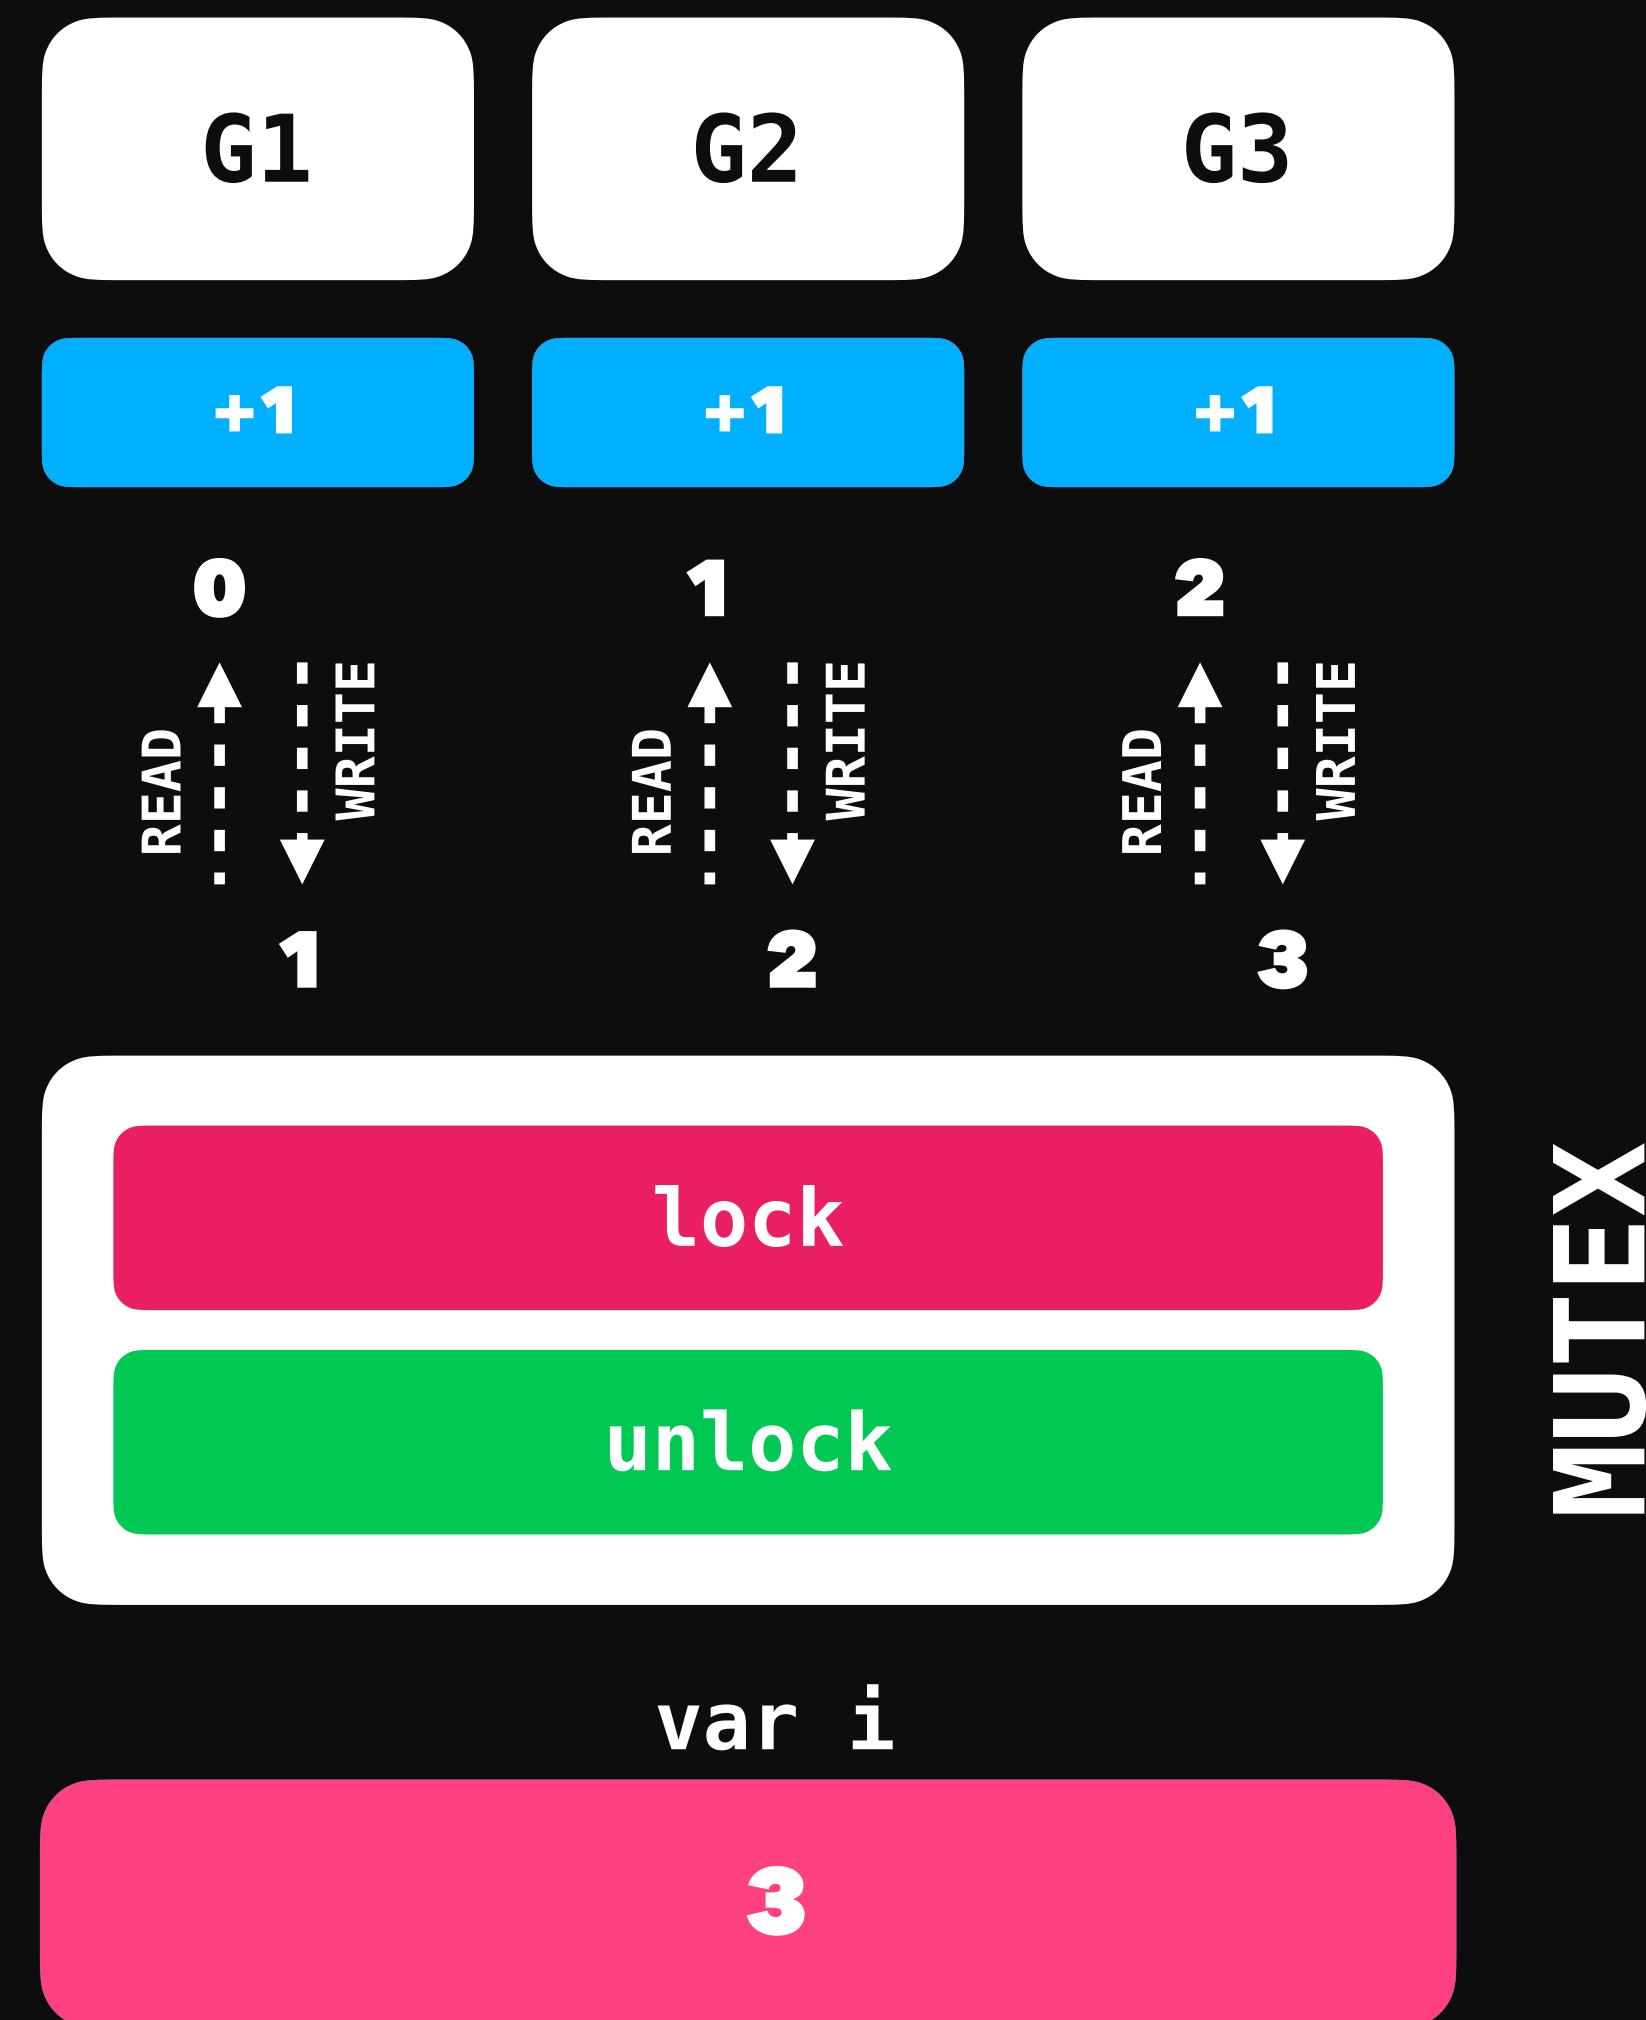
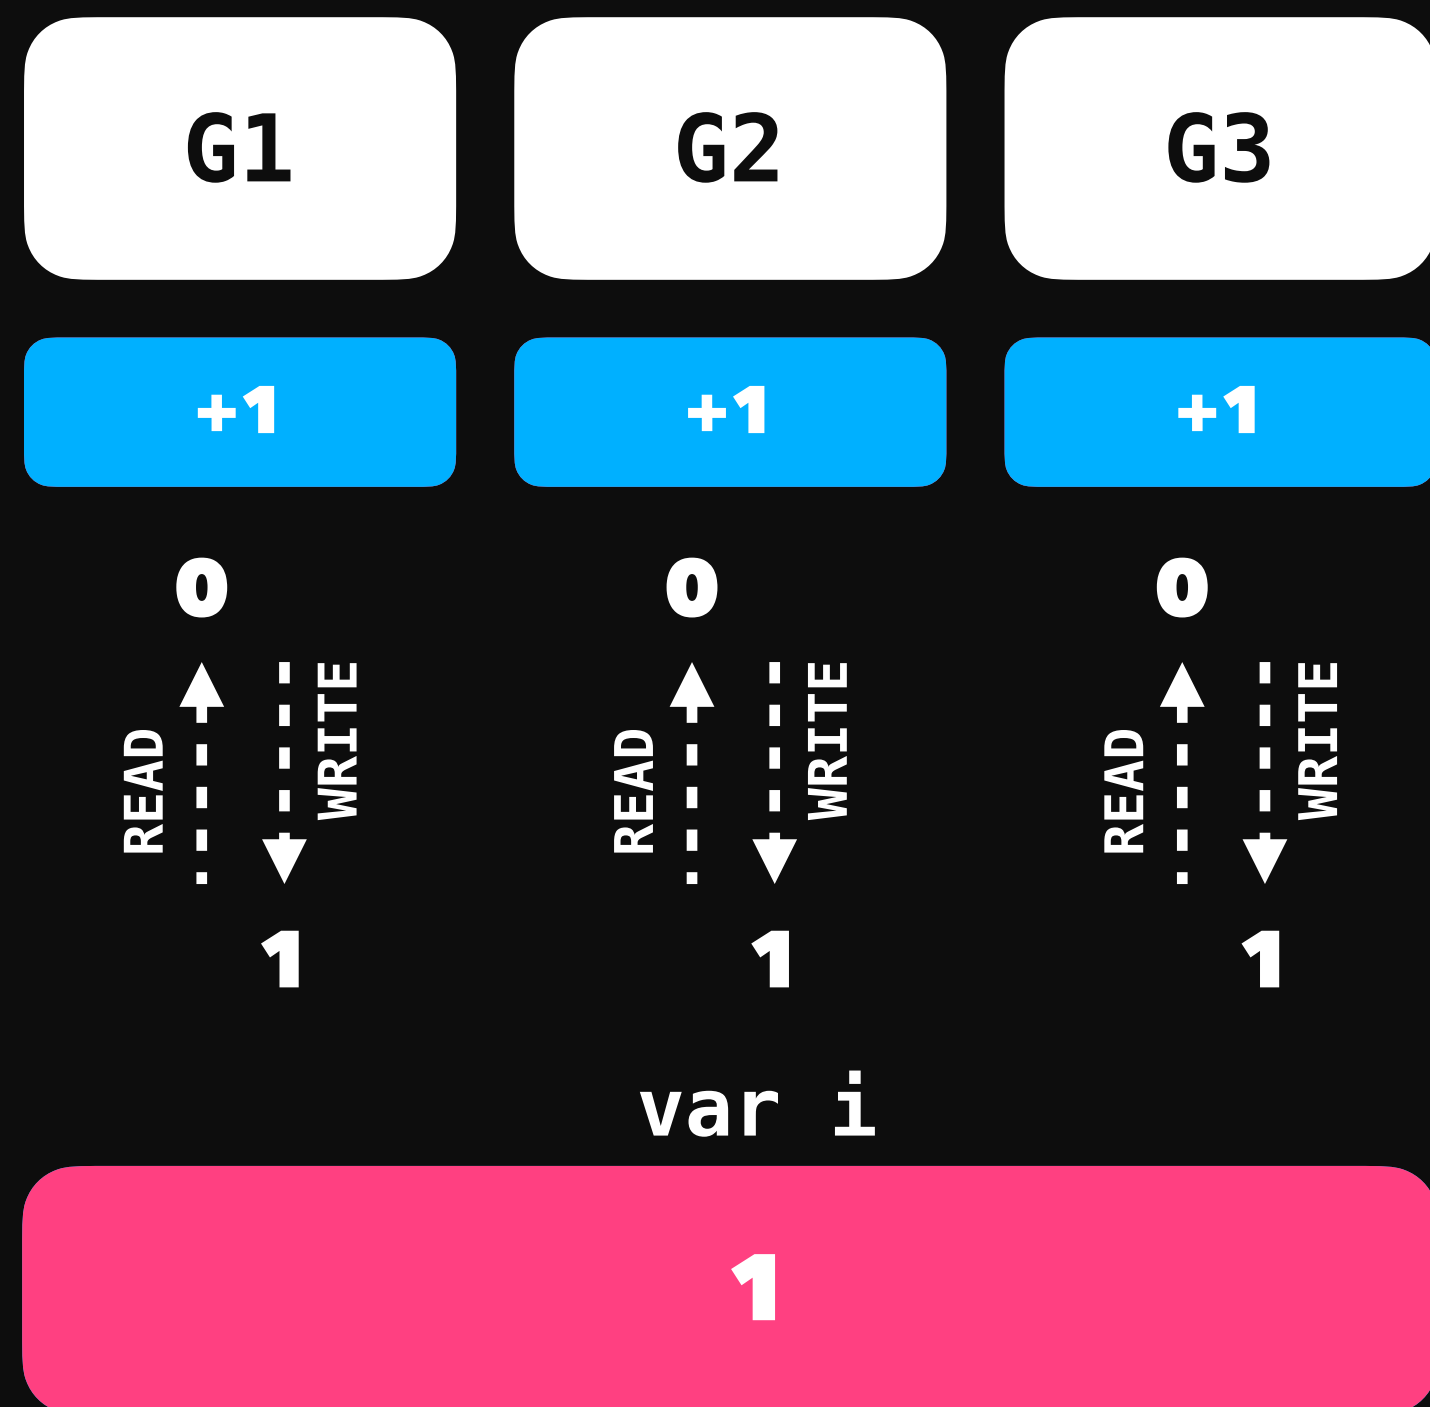
`increment value of i`


`store value of i`

INDIVISIBLE





UNINTERRUPTIBLE

ATOMIC





**IS THAT SO?**



var i






# Single Go Routine Context




NO MUTEX

# Main Go Routine Context



NEEDS MUTEX

# Multiple Go Routines Read Only Context



NO MUTEX

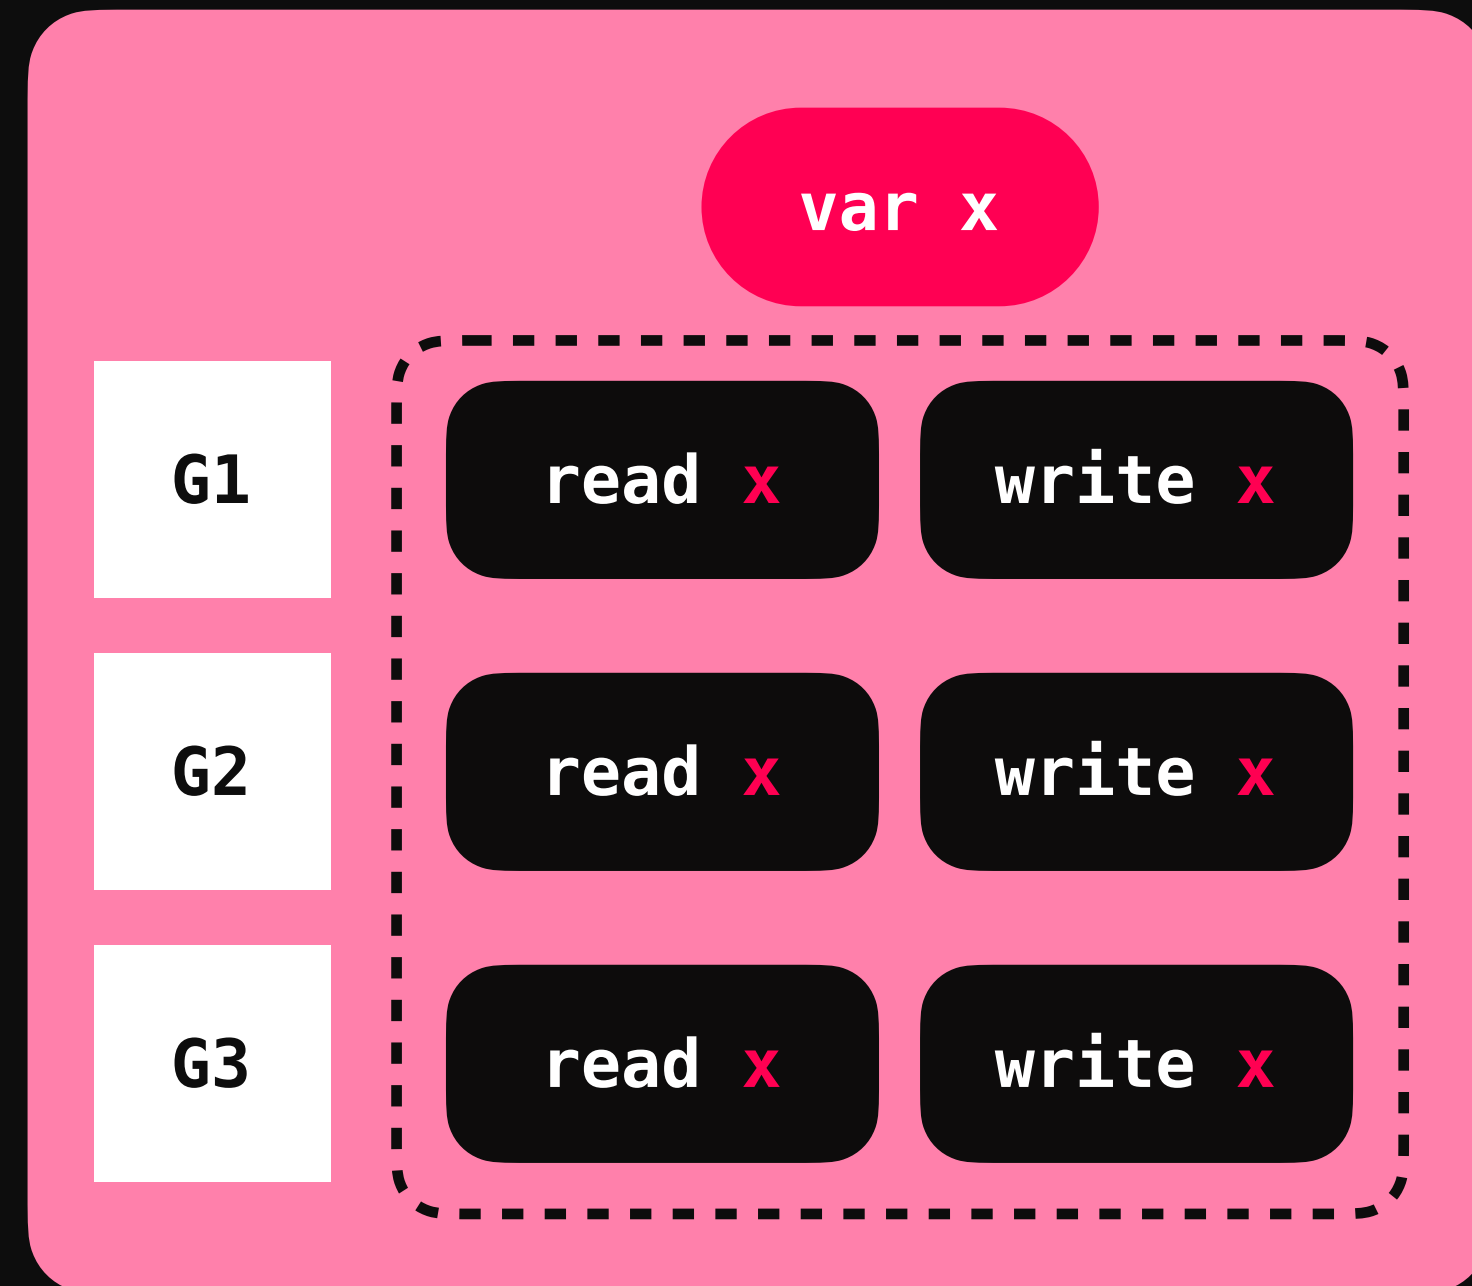
# Multiple Go Routines

## Read Write Context

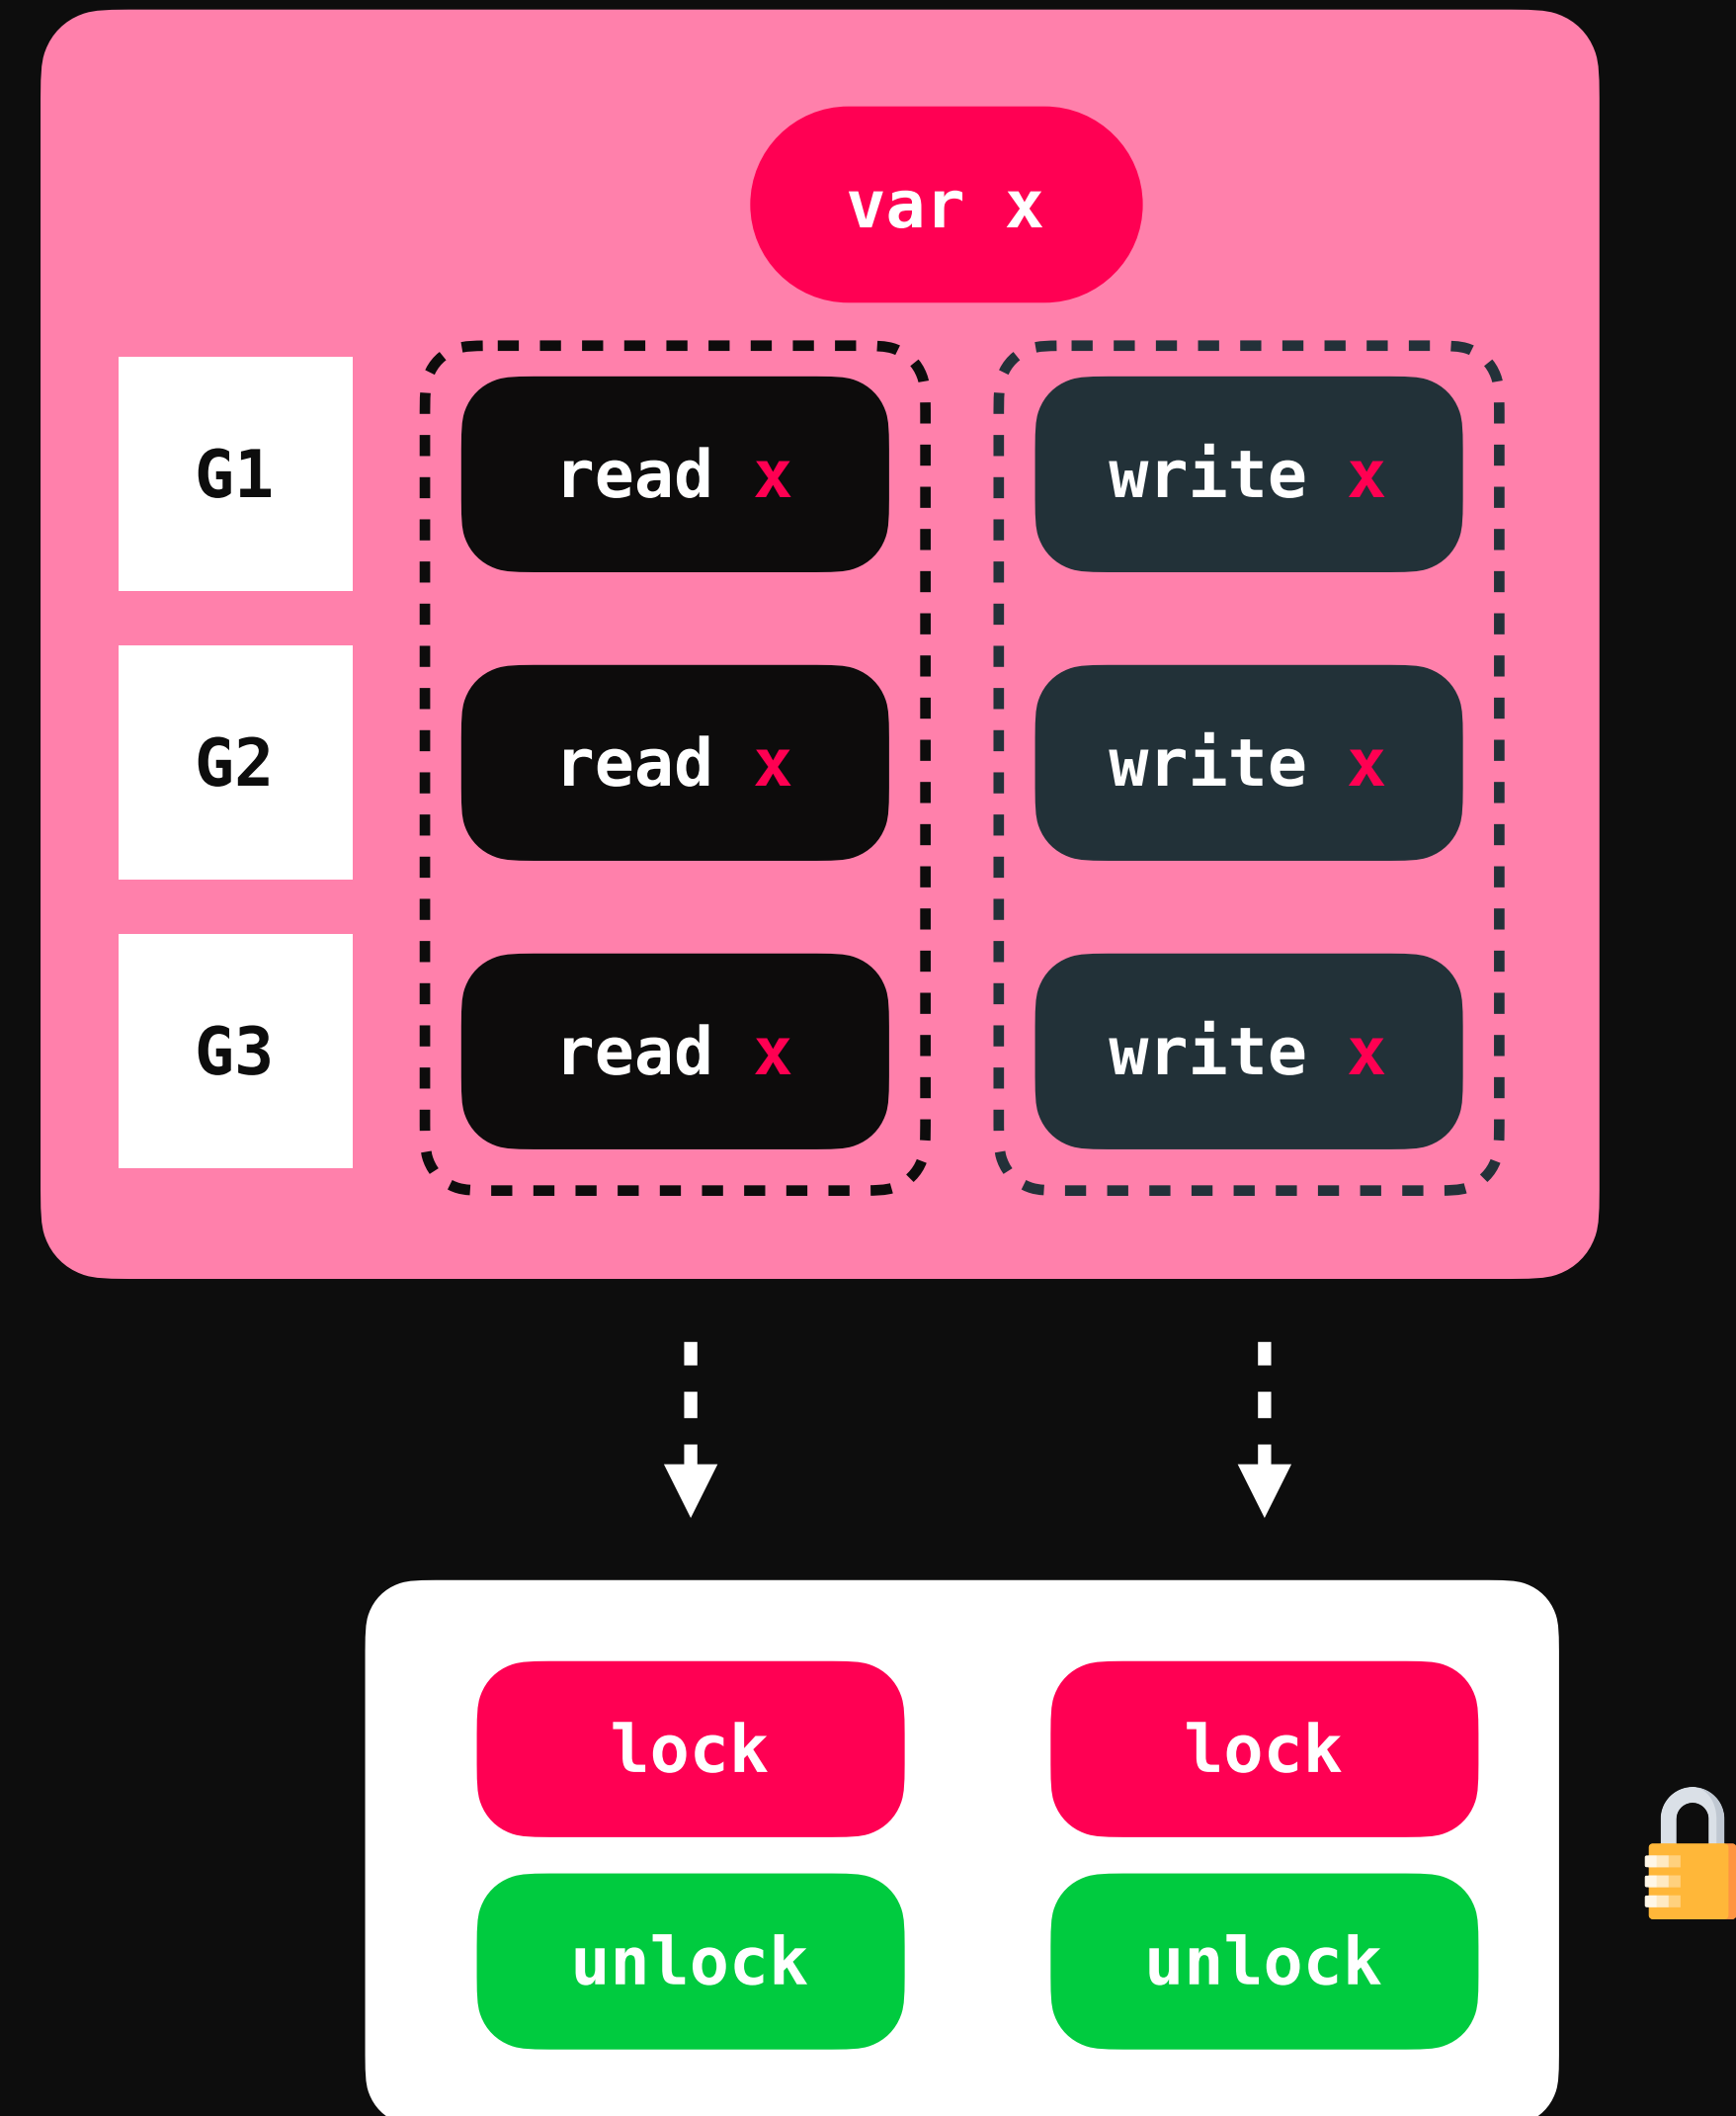


NEEDS MUTEX

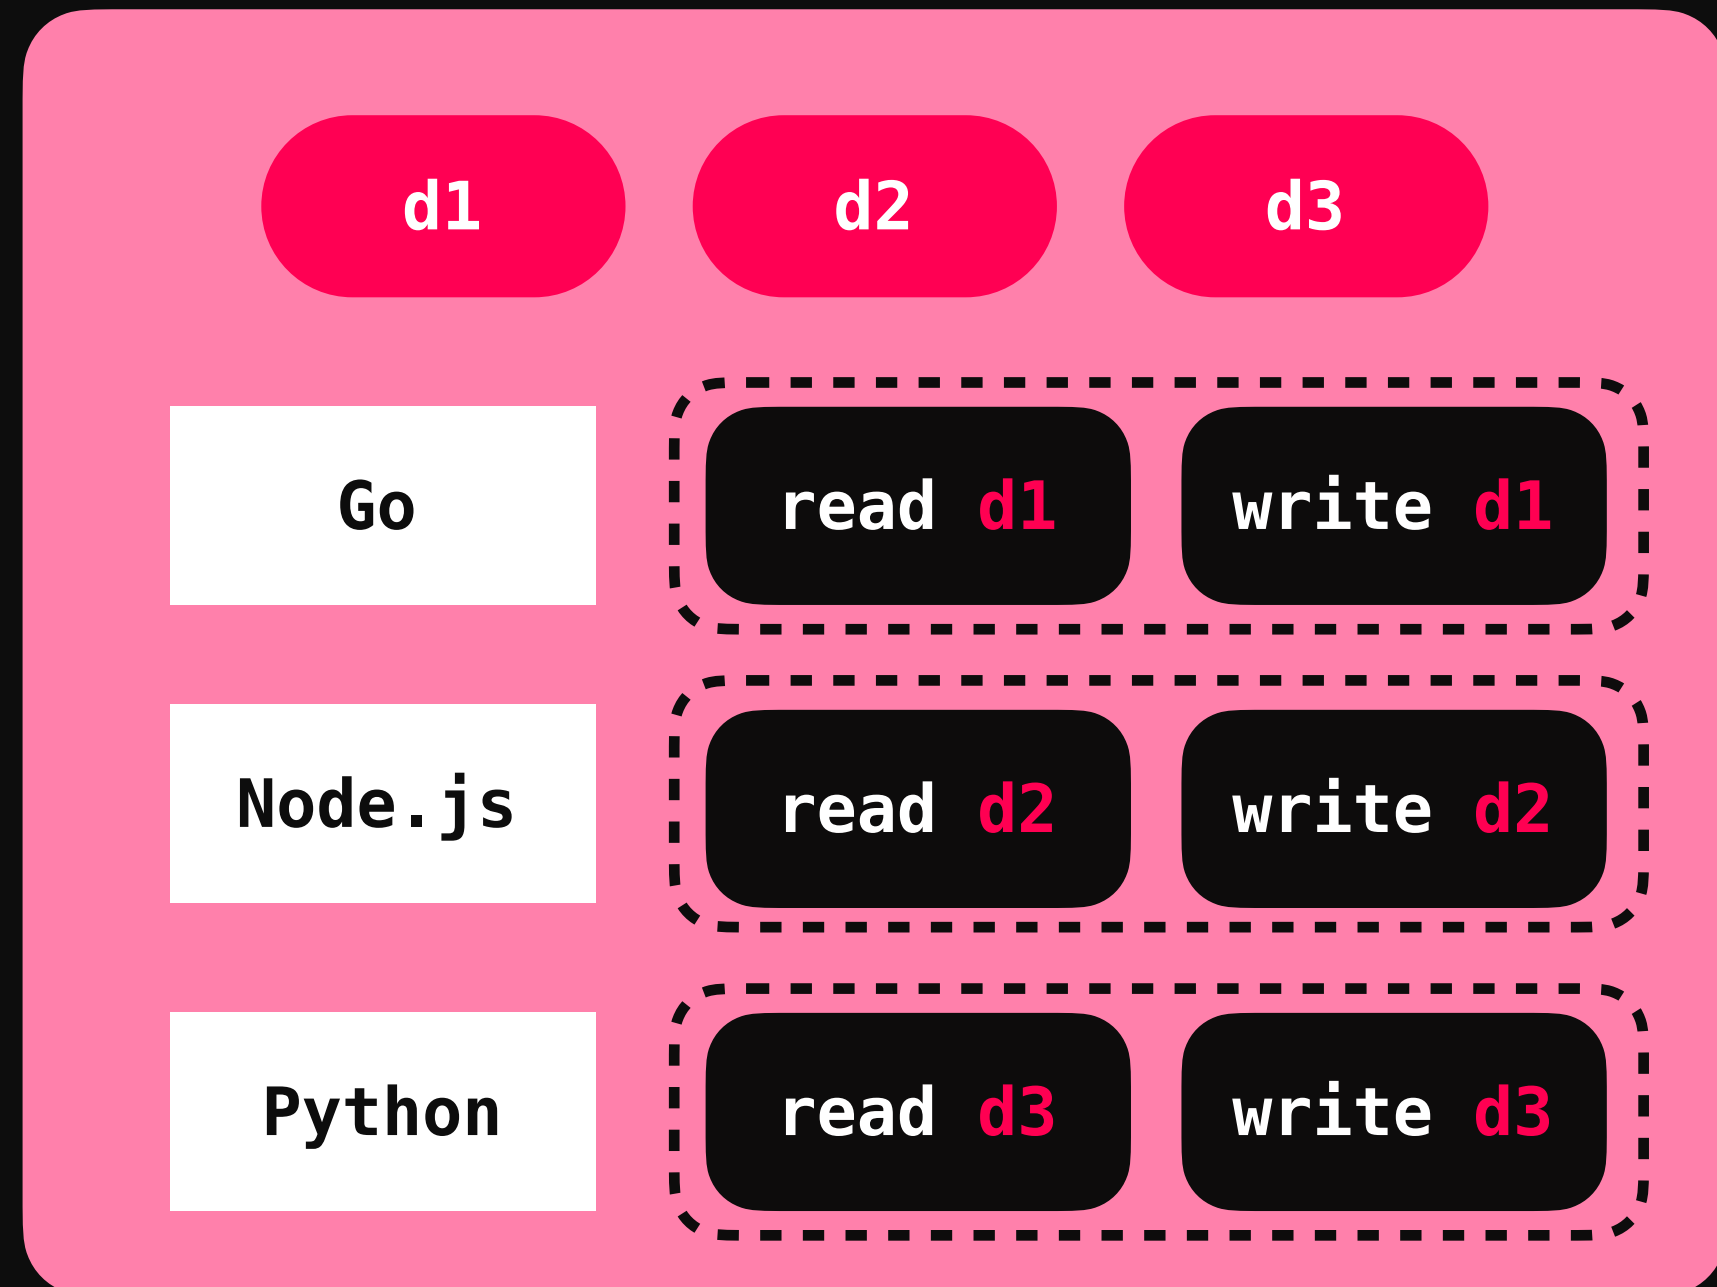
## Coarse Grained Context



## Fine Grained Context




## OS Context Different Locations





NO LOCK



## OS Context Same Location



NEEDS LOCK





MUTEX LOCK

⋮

⋮

G1

lock()

unlock()

G2

lock()

unlock()



DEADLOCK



# COFFMAN CONDITIONS



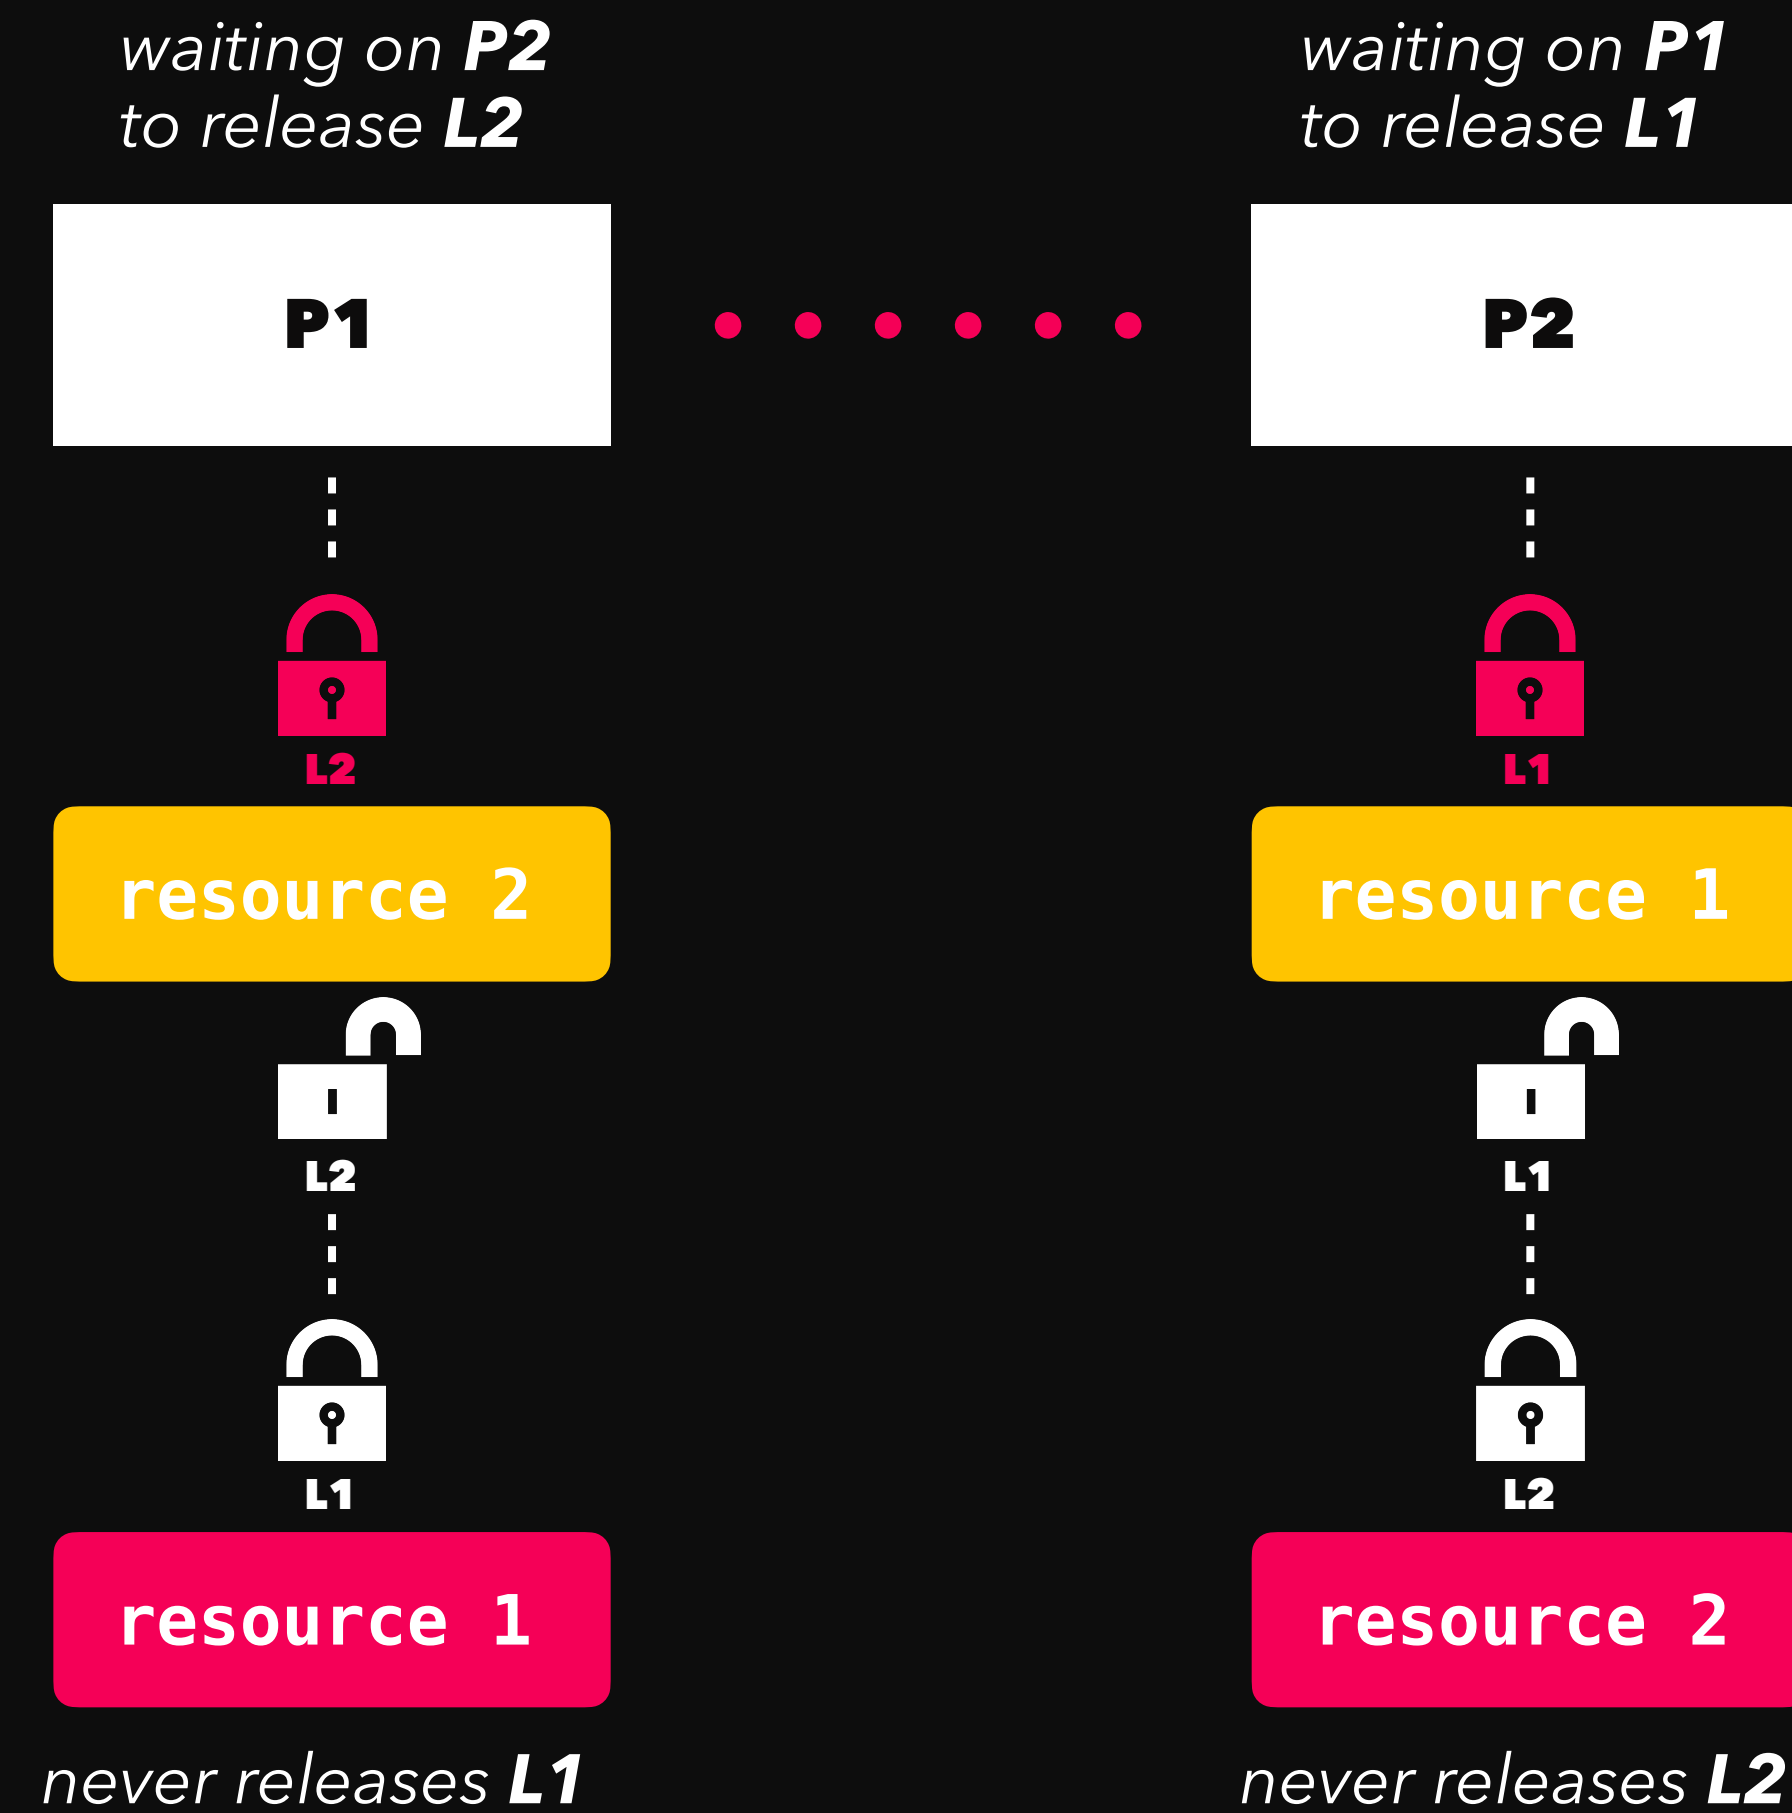
# WHAT IS A DEADLOCK?



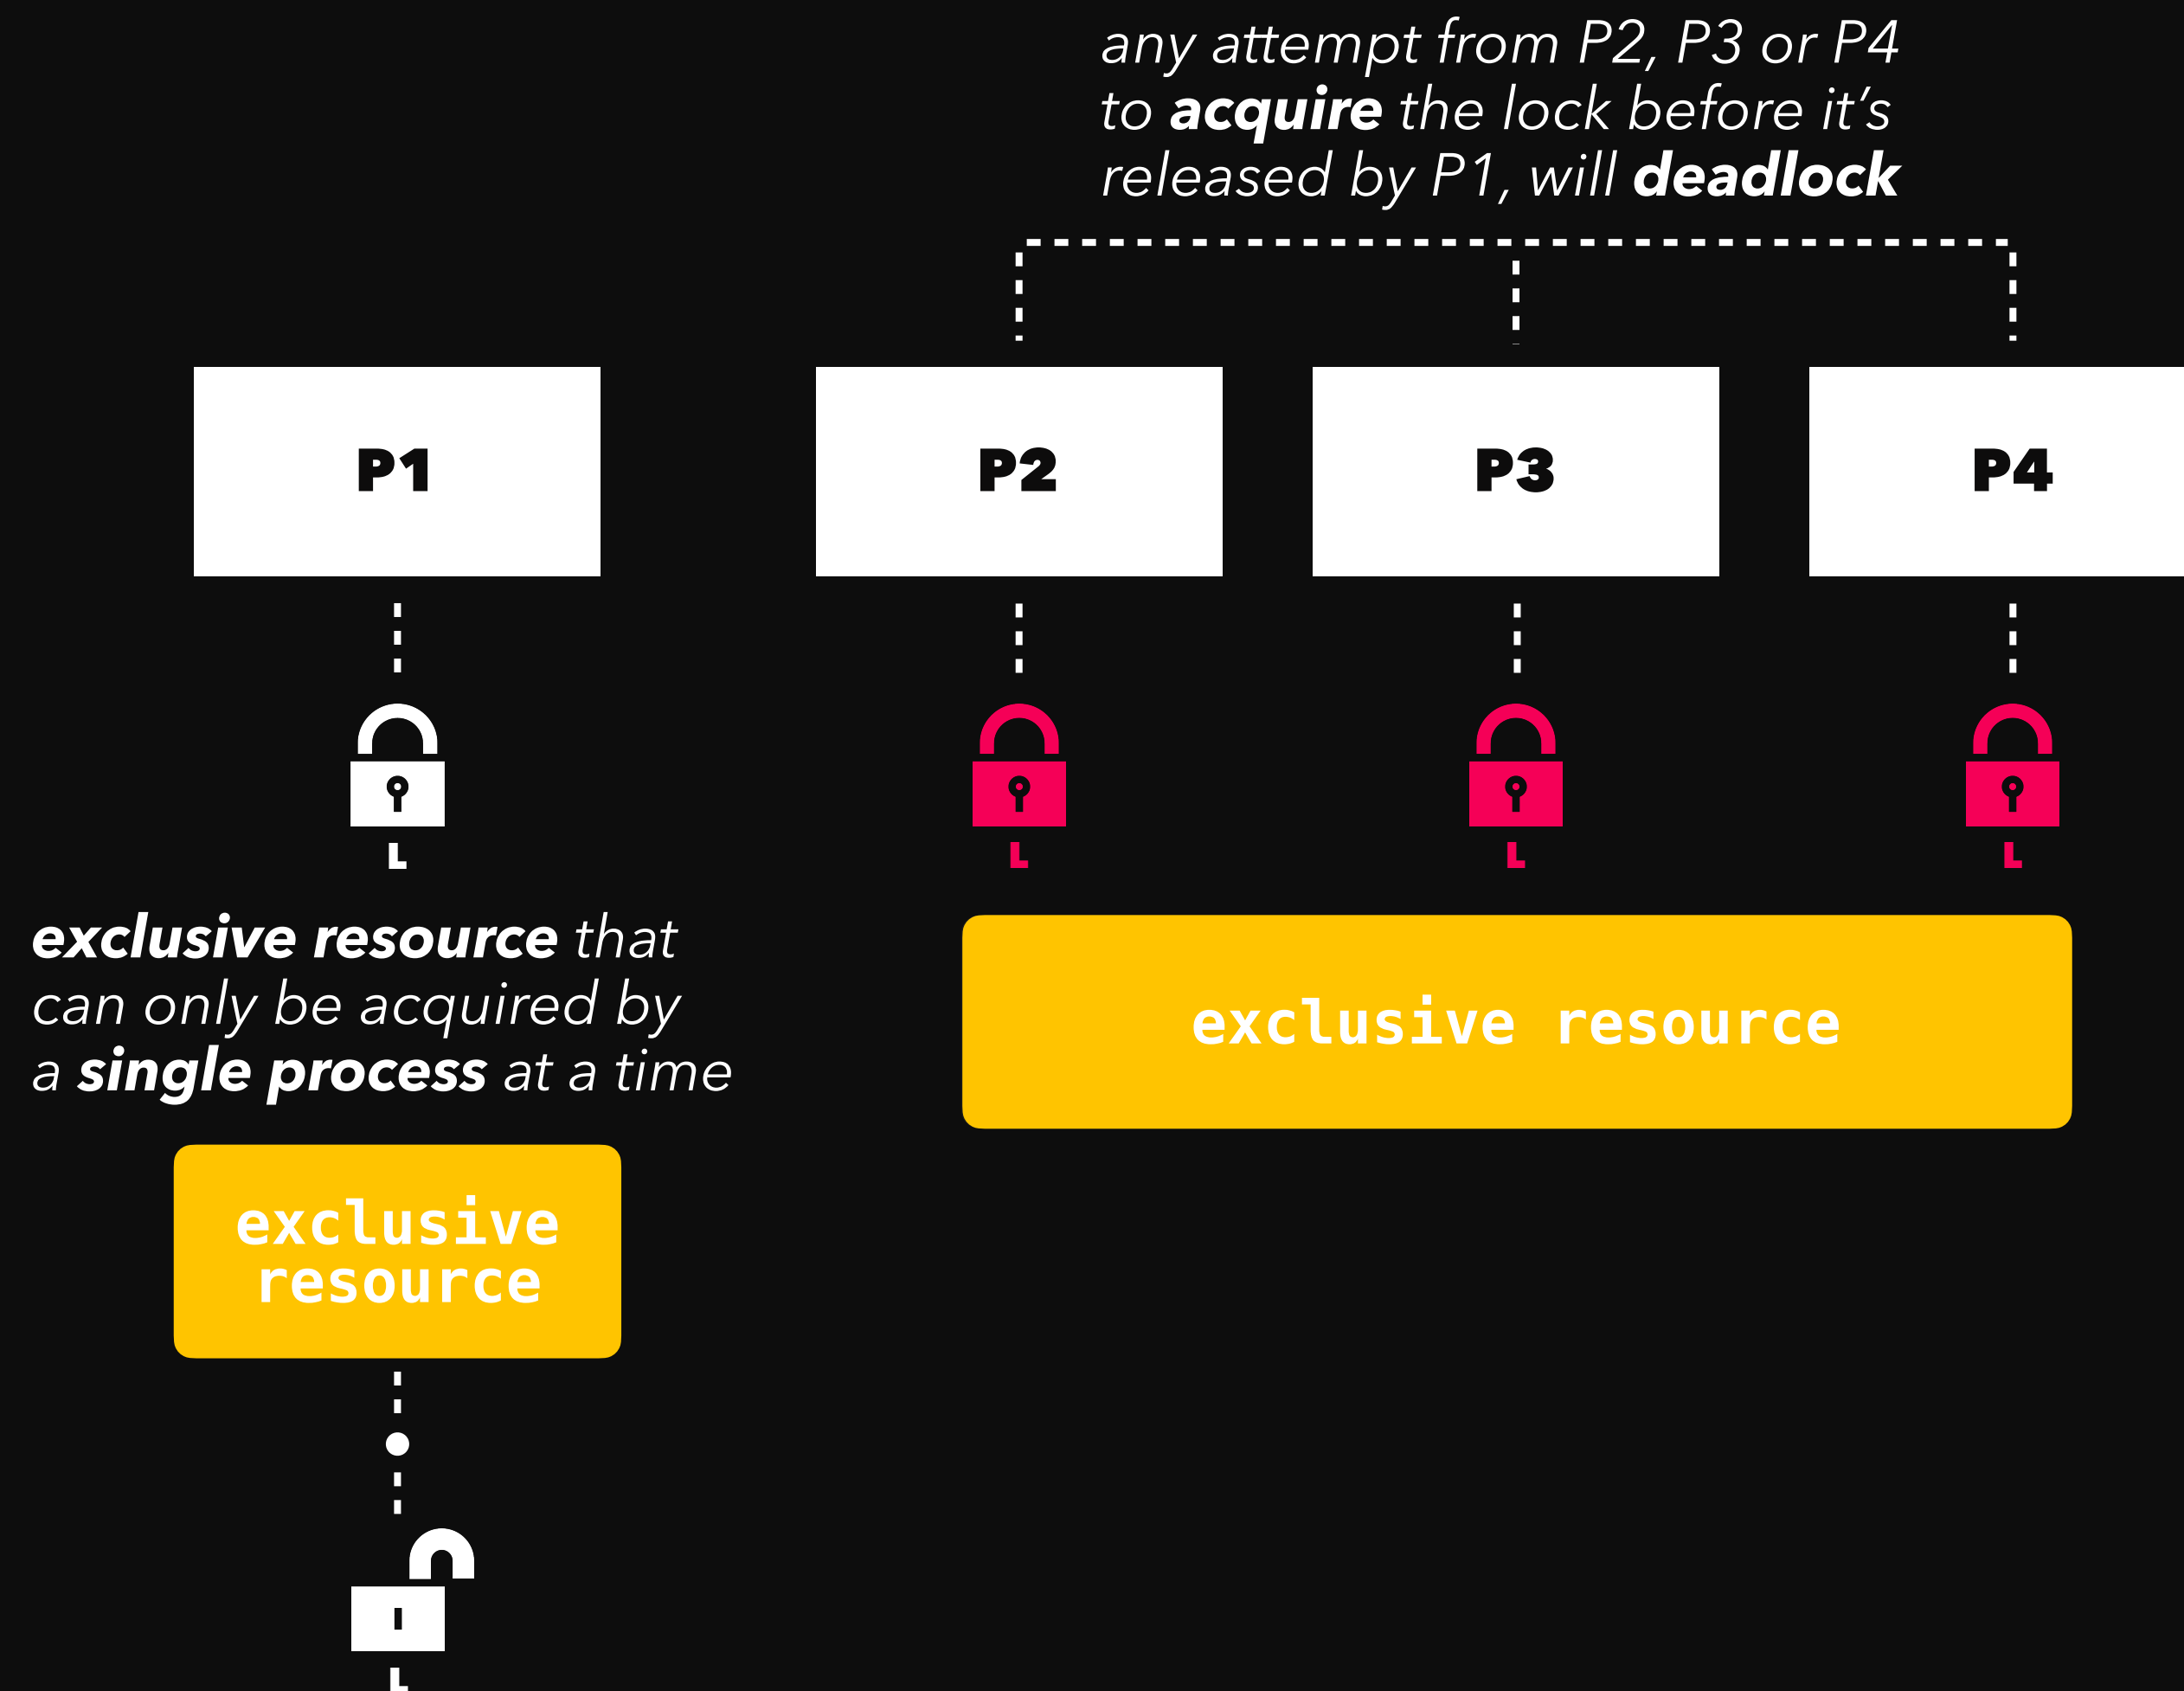
*A **deadlock** is a state in which each member of the group **waits** for another member, including itself, to take action, such as sending a message or more commonly **releasing a lock**.*



# 1 CIRCULAR WAIT




# 2 MUTUAL EXCLUSION



# 3 HOLD AND WAIT

*condition inside P1 is never satisfied*



*waiting on P1 to satisfy the condition*



# WHAT IS PREEMPTION?

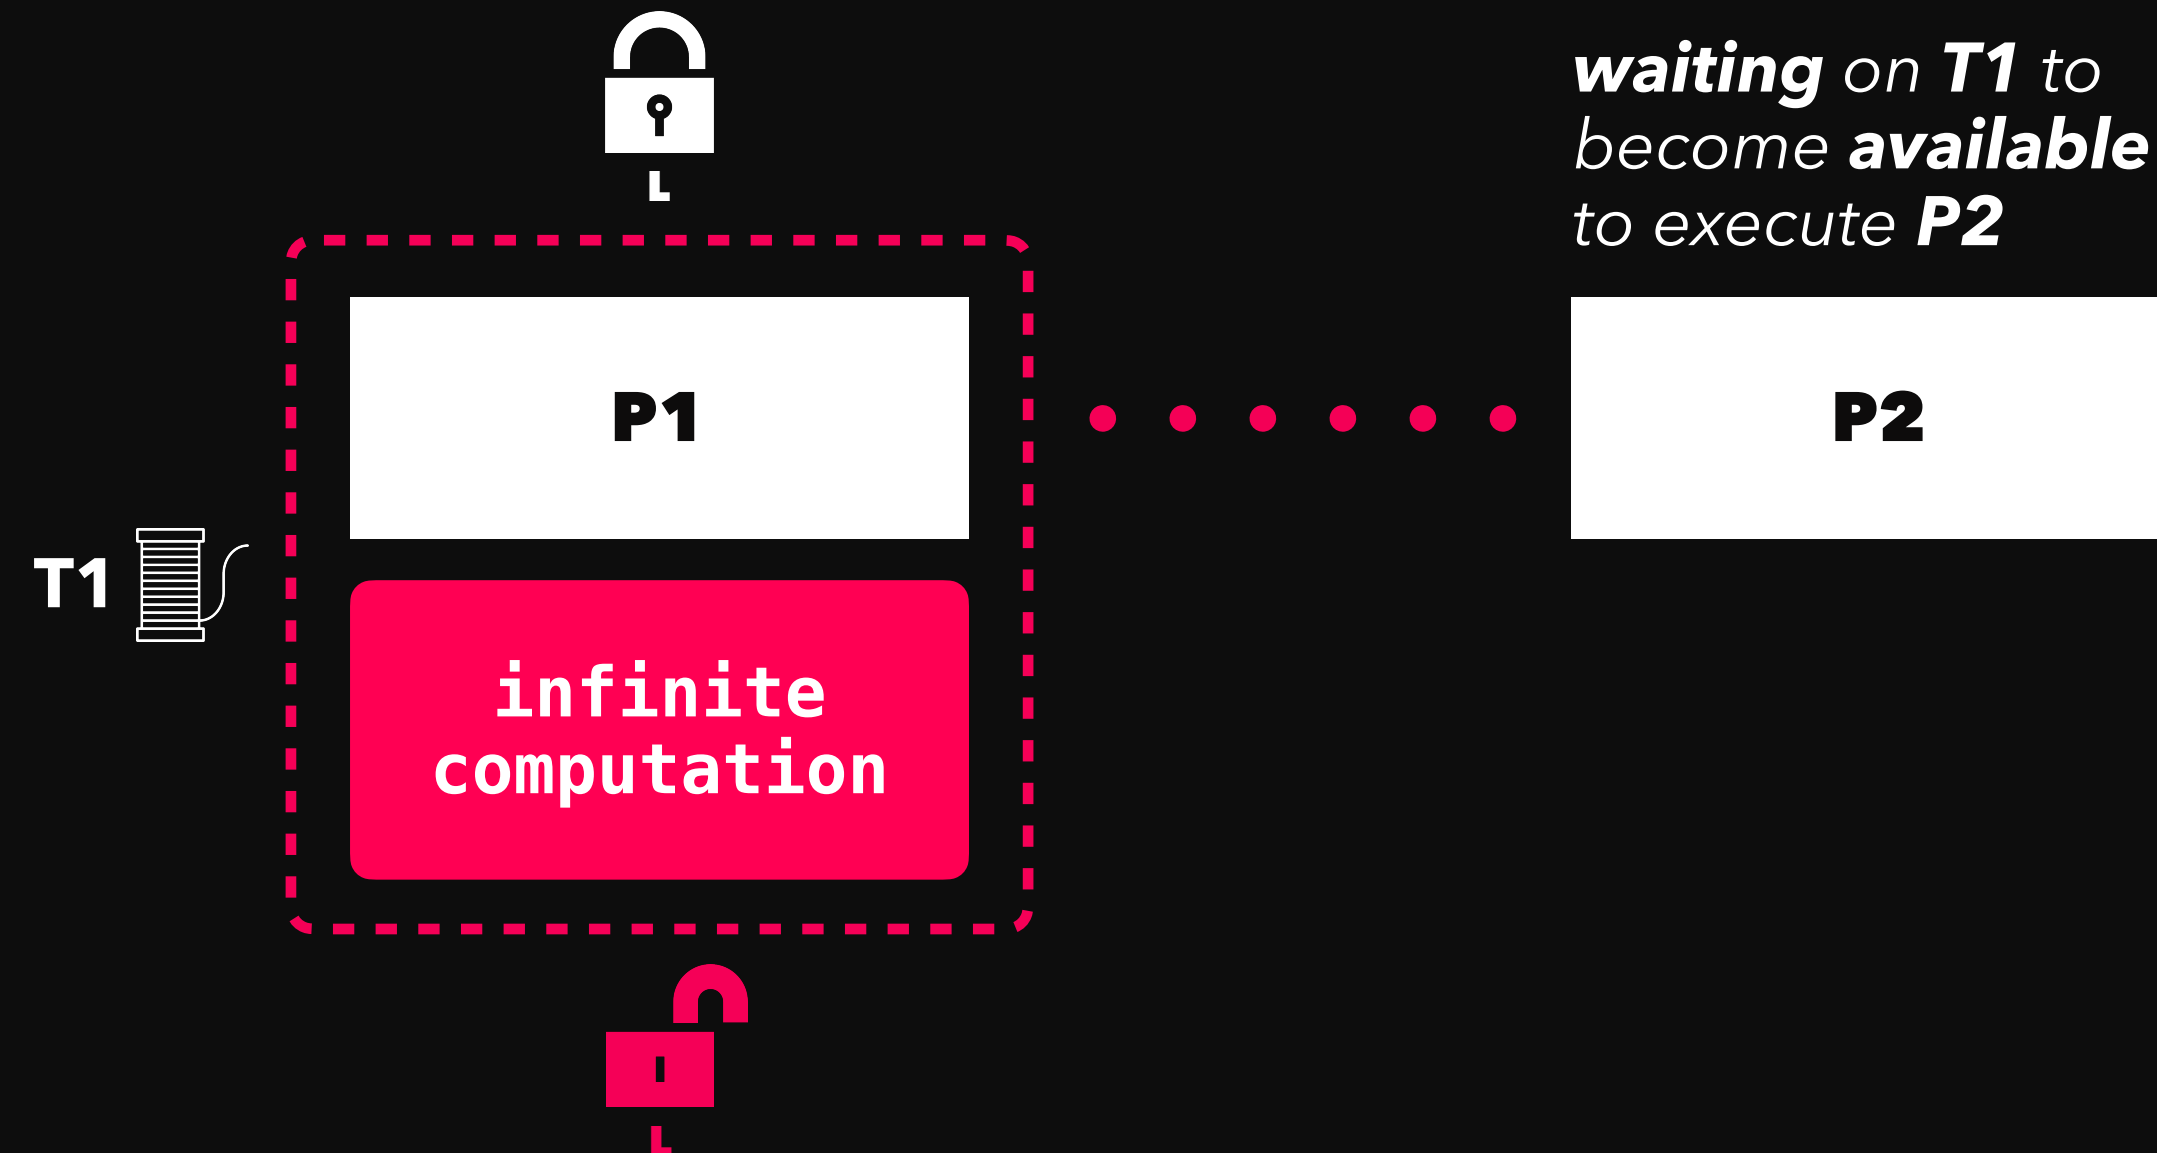


***Preemption** is the act of temporarily **interrupting** an **executing task**, with the intention of **resuming** it at a **later time**.*



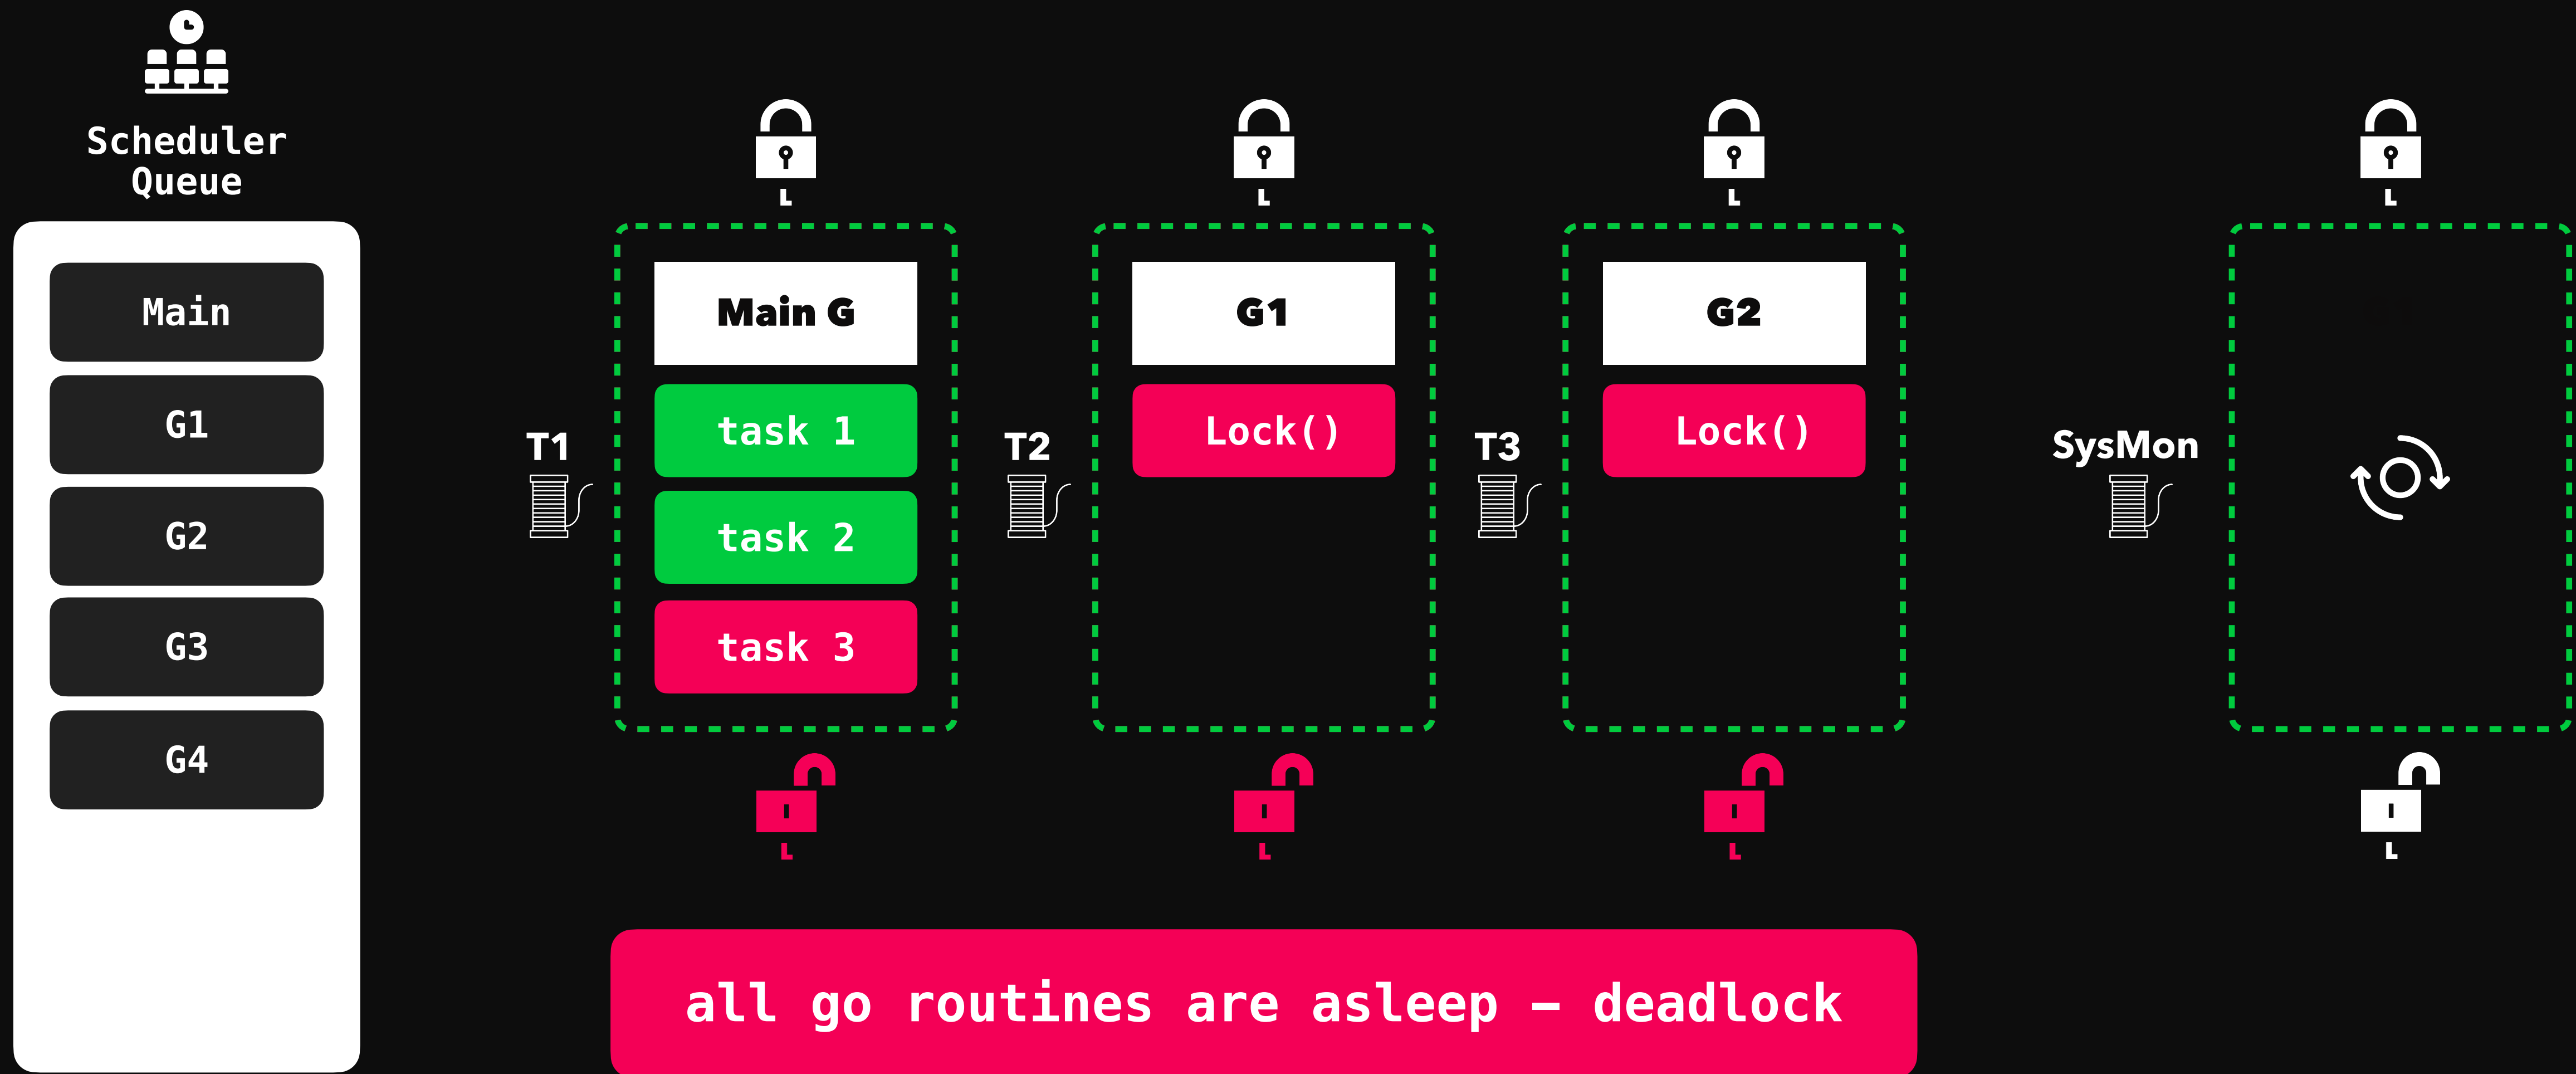
# 4 NO PREEMPTION

***T1** is always **stuck** executing **P1**, thus it's **never available** for other processes*



***waiting** on **T1** to become **available** to execute **P2***

# CHECK DEAD





Mike



17m

Jane





30m


John



15m



Mutex



released

G1

released

G2

G1 work

G2 work





**3**



**REFILL**



**1**



MUTEX LOCK

⋮

⋮

G1

acquire

work

sleep

release

100μs

G2

sleep

acquire

work

release

100μs



Scheduler Queue

G2



10ms

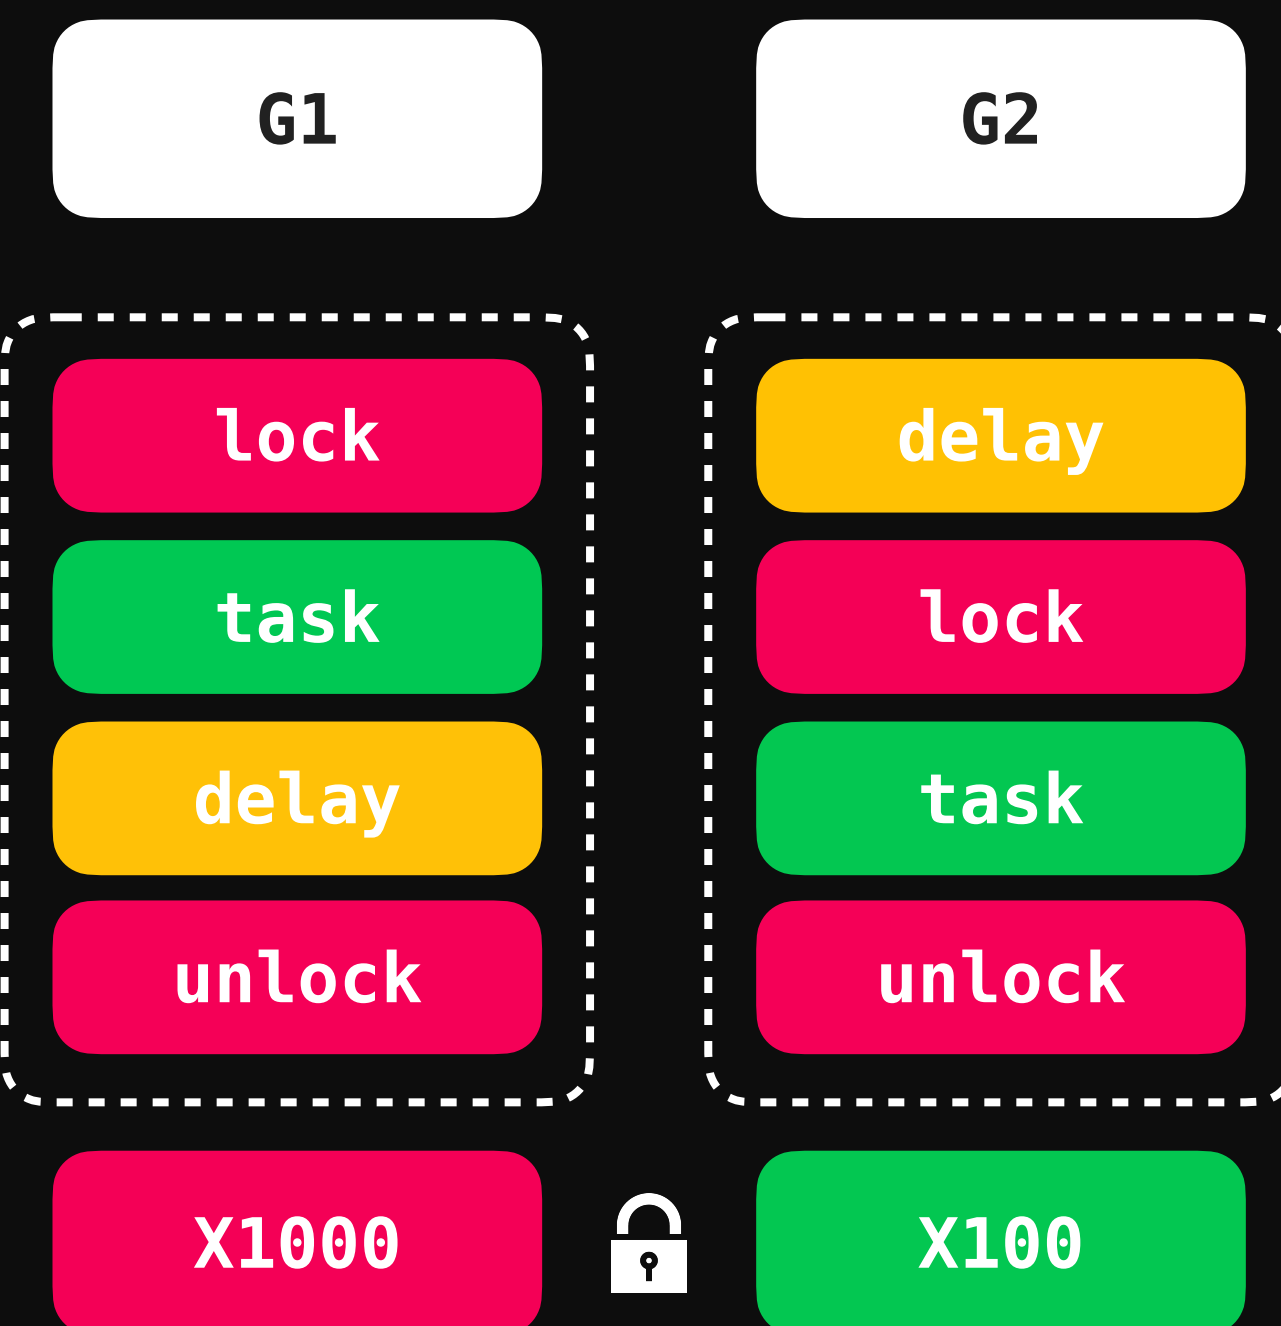
# CONTENTION




● UNEVEN WORK

VS

# STARVATION



● UNEVEN TIMING/HOLD TIME





MUTEX LOCK

G1

lock()

condition

unlock()

wait()

John



var left

var right

Hallway

G2

lock()

condition

unlock()

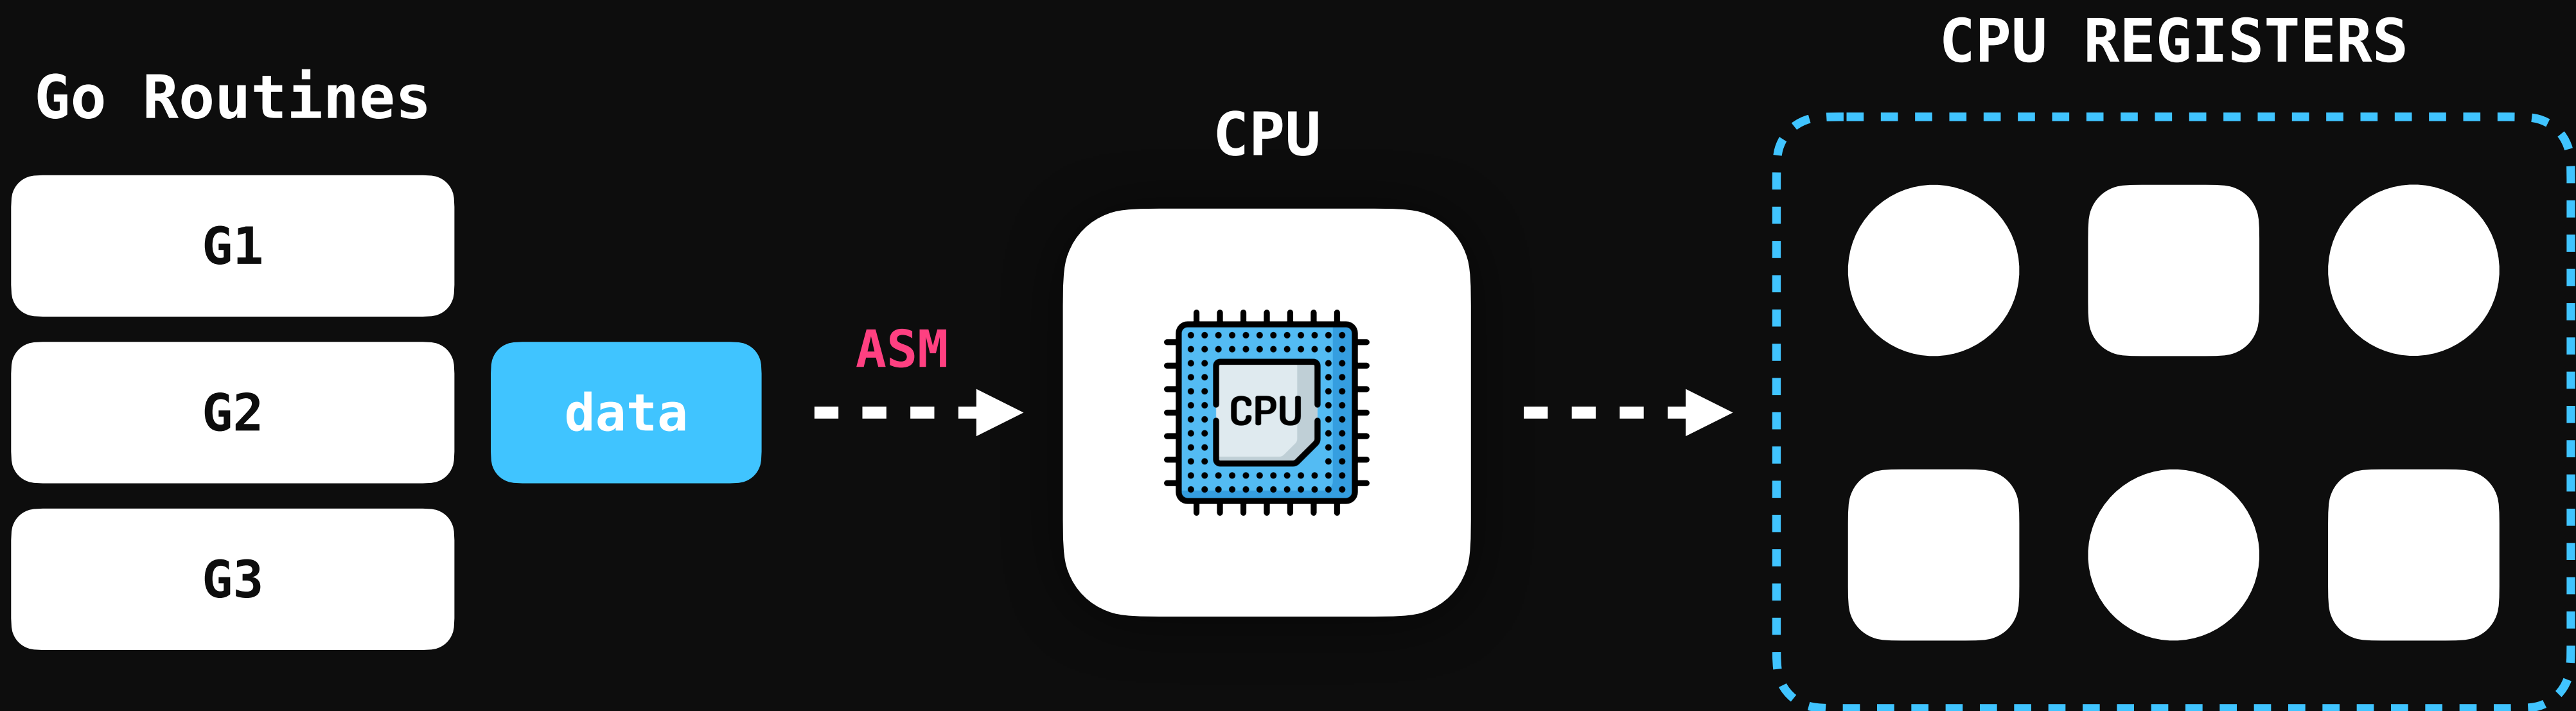
wait()


Alice

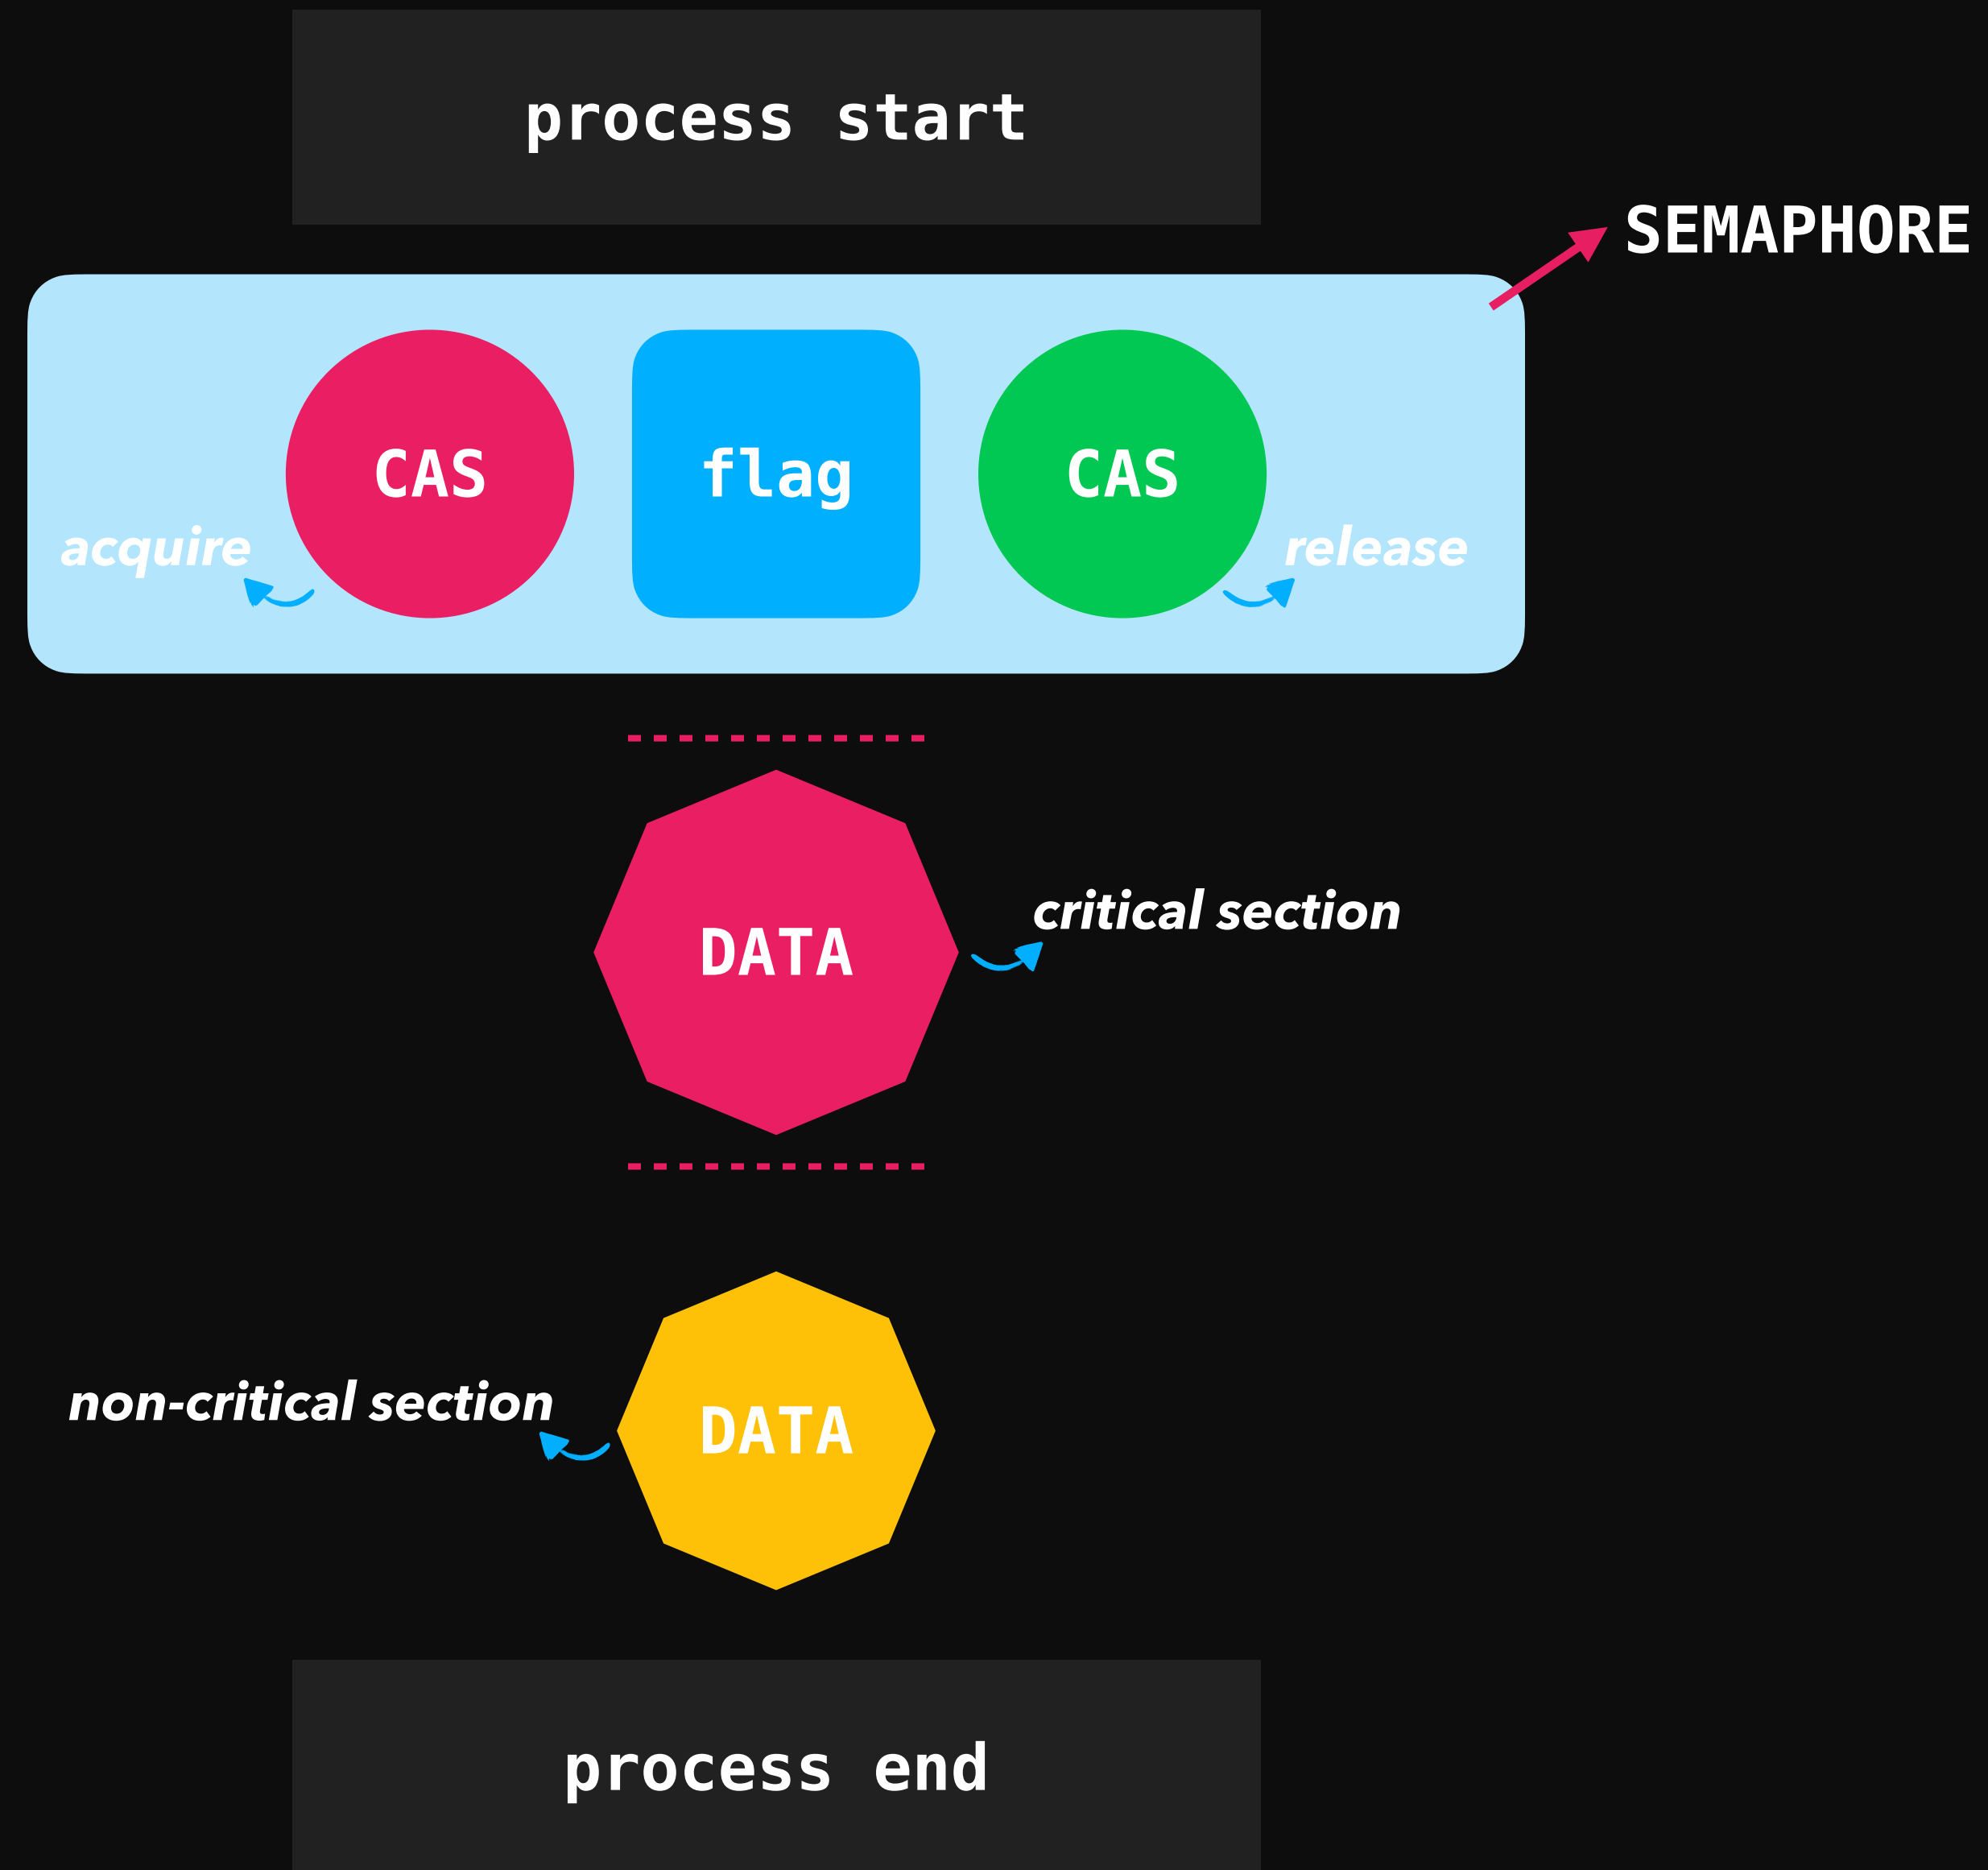
## sync.Locker

```
type Locker interface {  
    Lock()  
    Unlock()  
}
```











*Lock()*  
*Unlock()*

**TEST & SET**

Prefer **Atomics** over **Mutexes** for **simple data**

Use **Mutex** for write heavy scenarios

Use **RWMutex** for read heavy / mixed scenarios

Prefer **Fine Grained Context** where possible (limit the context)

**Don't use** the mutex **longer** than you need to (**avoid extra hold time**)


Use **types**, a **local mutex** and **methods** over direct mutex calls

Avoid **Contention** by **distributing work evenly**

Avoid **Starvation** by testing for **mutex fairness**

Use **sync.Locker** for generic code that uses a mutex

**Concurrency Control** ensures that **correct results** for **concurrent operations** are generated, while getting those results as **quickly** as possible.



# ACID

**ATOMICITY**

**CONSISTENCY**

**ISOLATION**

**DURABILITY**

# TRANSACTION








\$100  
-\$40



Alice



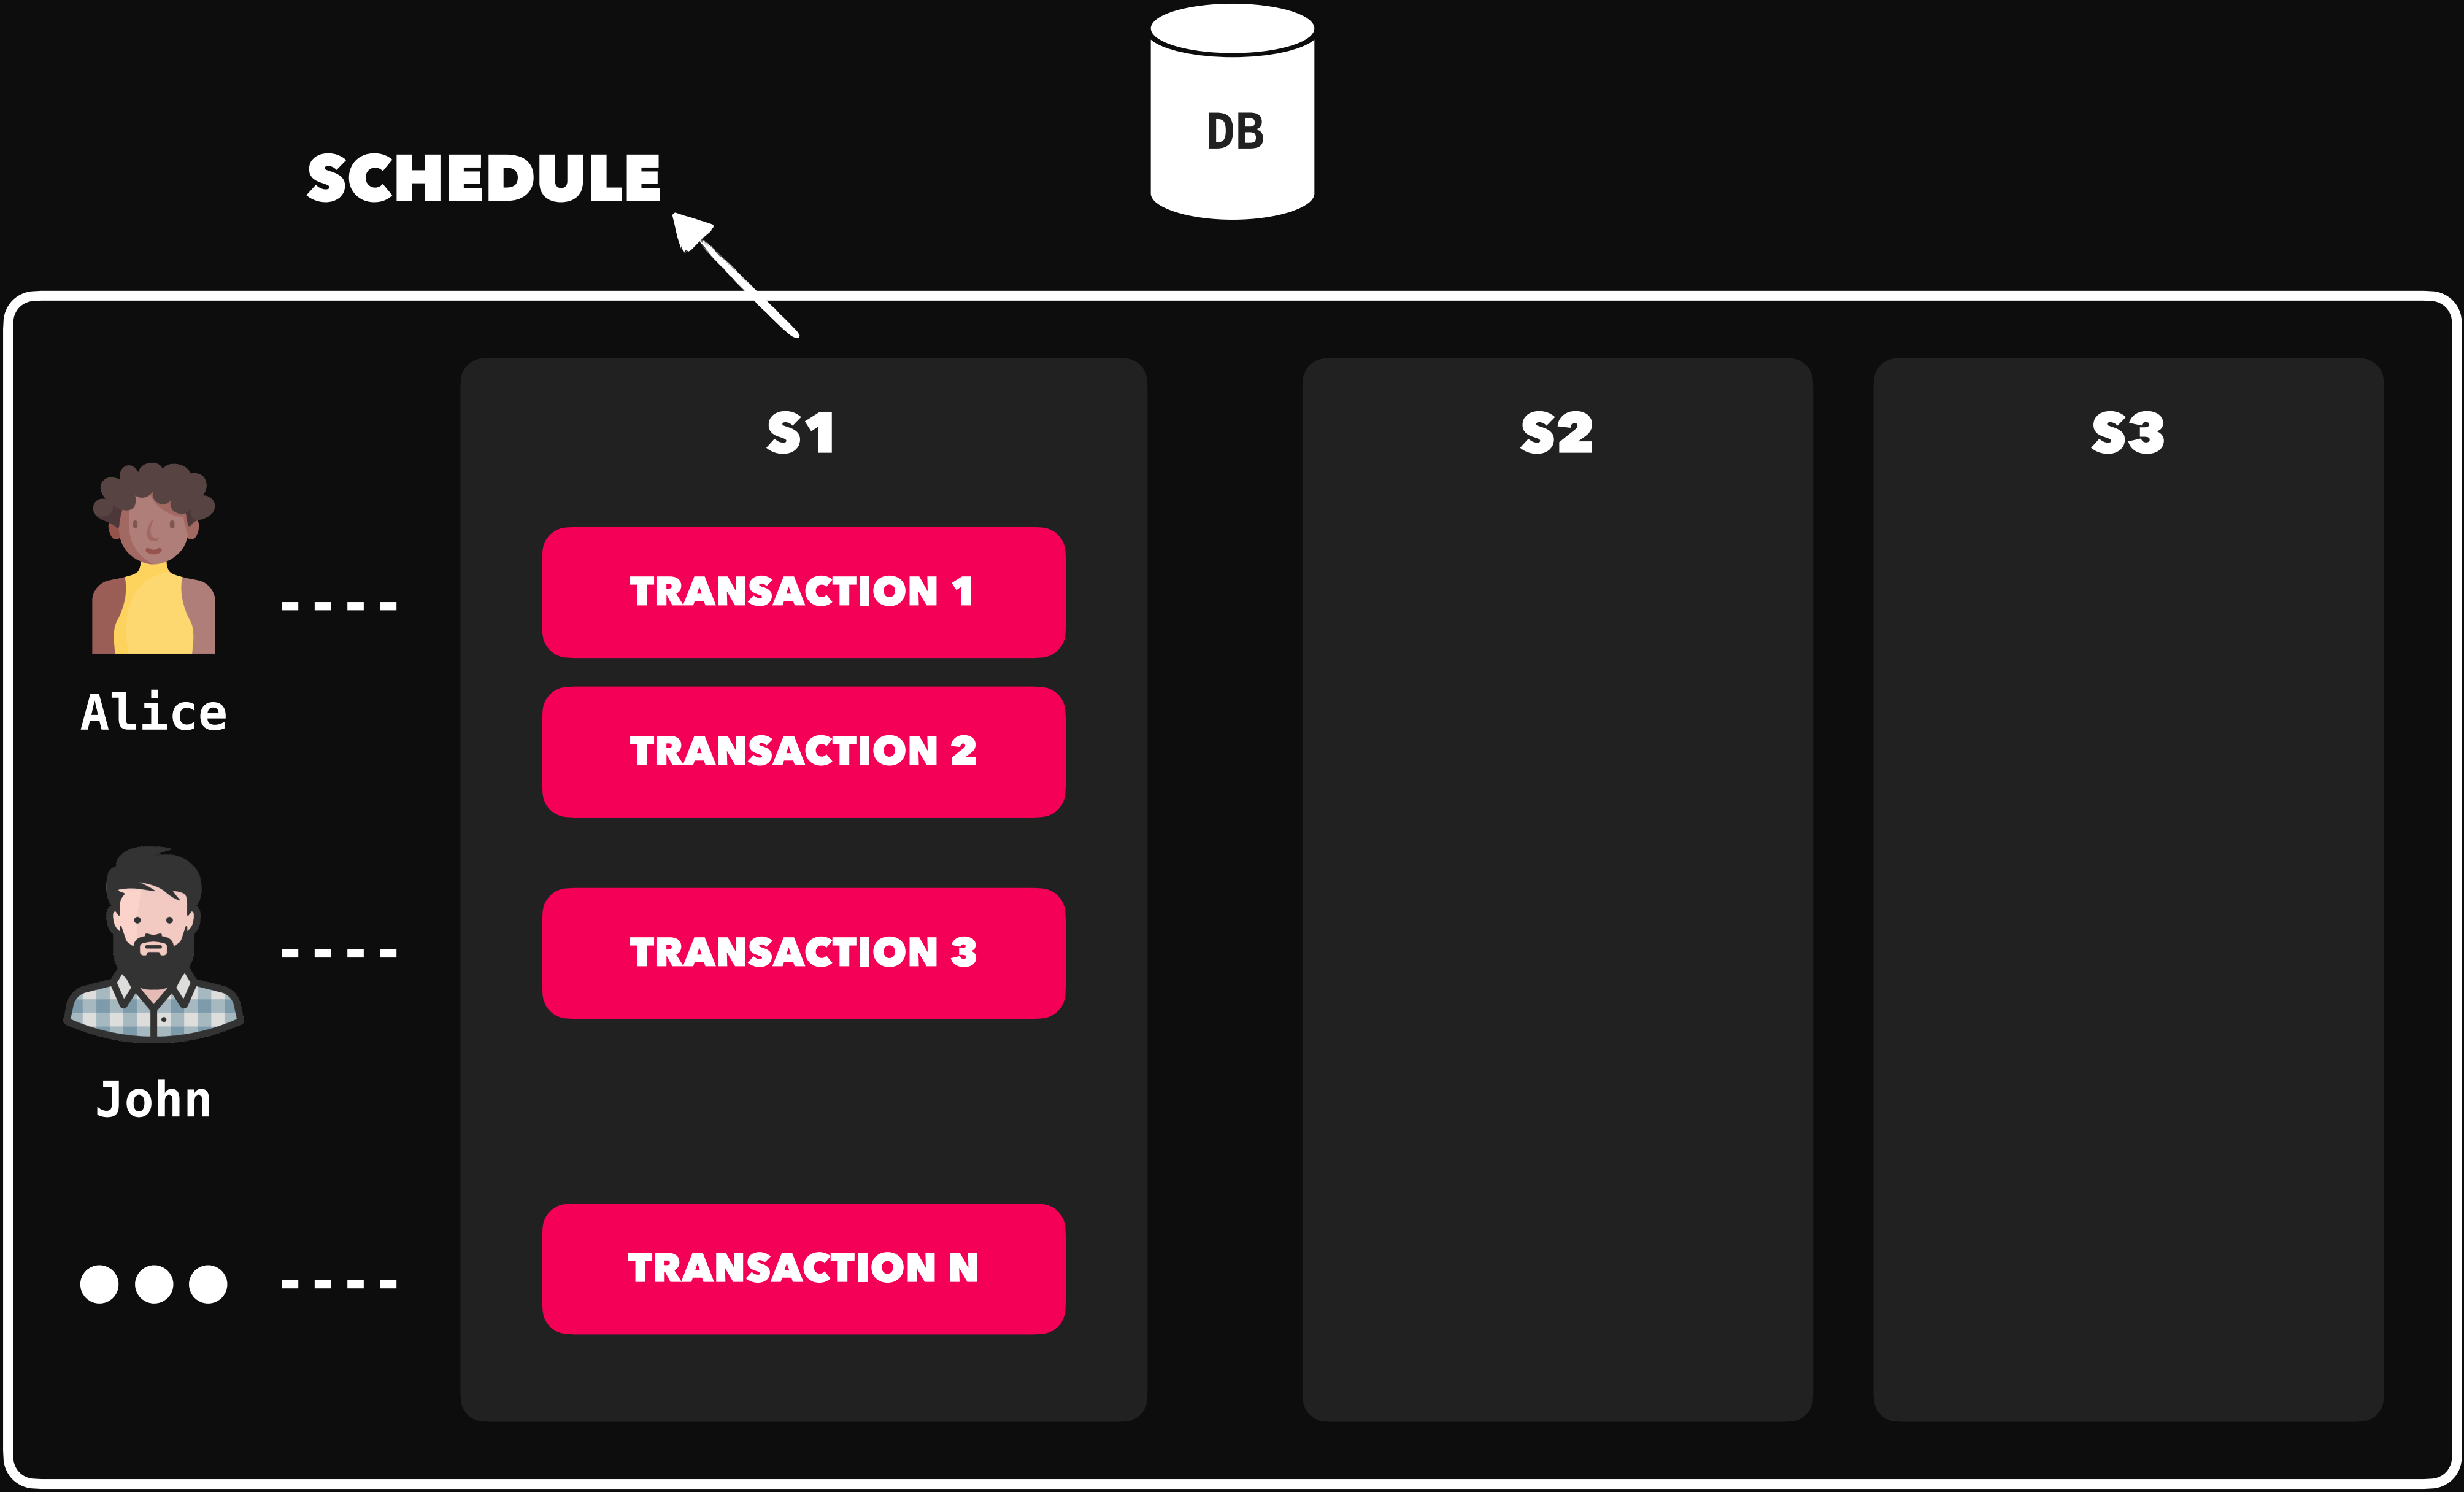
\$0  
+\$40



Book Store

No **payment**  
without the **book**

No **book** without  
**successful payment**



Alice wants to initiate  
2 **transactions** at the  
**same time**, transferring  
John **\$100/transaction**



**\$200**



**\$150**  
**-\$200**



**\$0**  
**+\$200**



**Alice**

**TX1**

**\$100**





**TX2**

**\$100**



**John**

Alice wants to initiate  
2 **transactions** at the  
**same time**, transferring  
John **\$100/transaction**



Alice wants to initiate  
2 **transactions** at the  
**same time**, transferring  
John **\$100/transaction**



\$150




\$150  
-\$100



\$0  
+\$100




Alice




John

1d

In **Concurrency Control** of databases and various **transactional applications**, a transaction schedule is **Serializable** if its **outcome** is **equal** to the outcome of its transactions **executed serially**, without overlapping in time.



# SCHEDULING




**SERIAL**

**T1**


**T2**

**Serial Schedule**



**CORRECT**, BUT SLOW

**NON-SERIAL**



**T1**

**T2**


**Non-Serial Schedule**



**CORRECT & FAST**  
**SERIALIZABLE**



**FAST**, BUT INCORRECT

# SERIALIZABILITY



**RESULT EQUIVALENT**

**CONFLICT SERIALIZABLE**



**VIEW SERIALIZABLE**



# RESULT EQUIVALENT

Initial values

X: 2  
Y: 5

S1

S2



Initial values

X: 3  
Y: 6

S1

S2



T1

R(X)

X=X+5

W(X)

R(Y)

Y=Y+5

W(Y)

T2

R(X)

X=X\*3

W(X)

X: 21

Y: 10

T1

R(X)

X=X+5

W(X)

T2

R(X)

X=X\*3

W(X)

T1

R(Y)

Y=Y+5

W(Y)

X: 21

Y: 10

T1

R(X)

X=X+1

W(X)

R(Y)

Y=Y\*2

W(Y)

T2

R(X)

X=X\*3

W(X)

X: 12

Y: 12

T1

R(X)

X=X+1

W(X)

T2

R(Y)

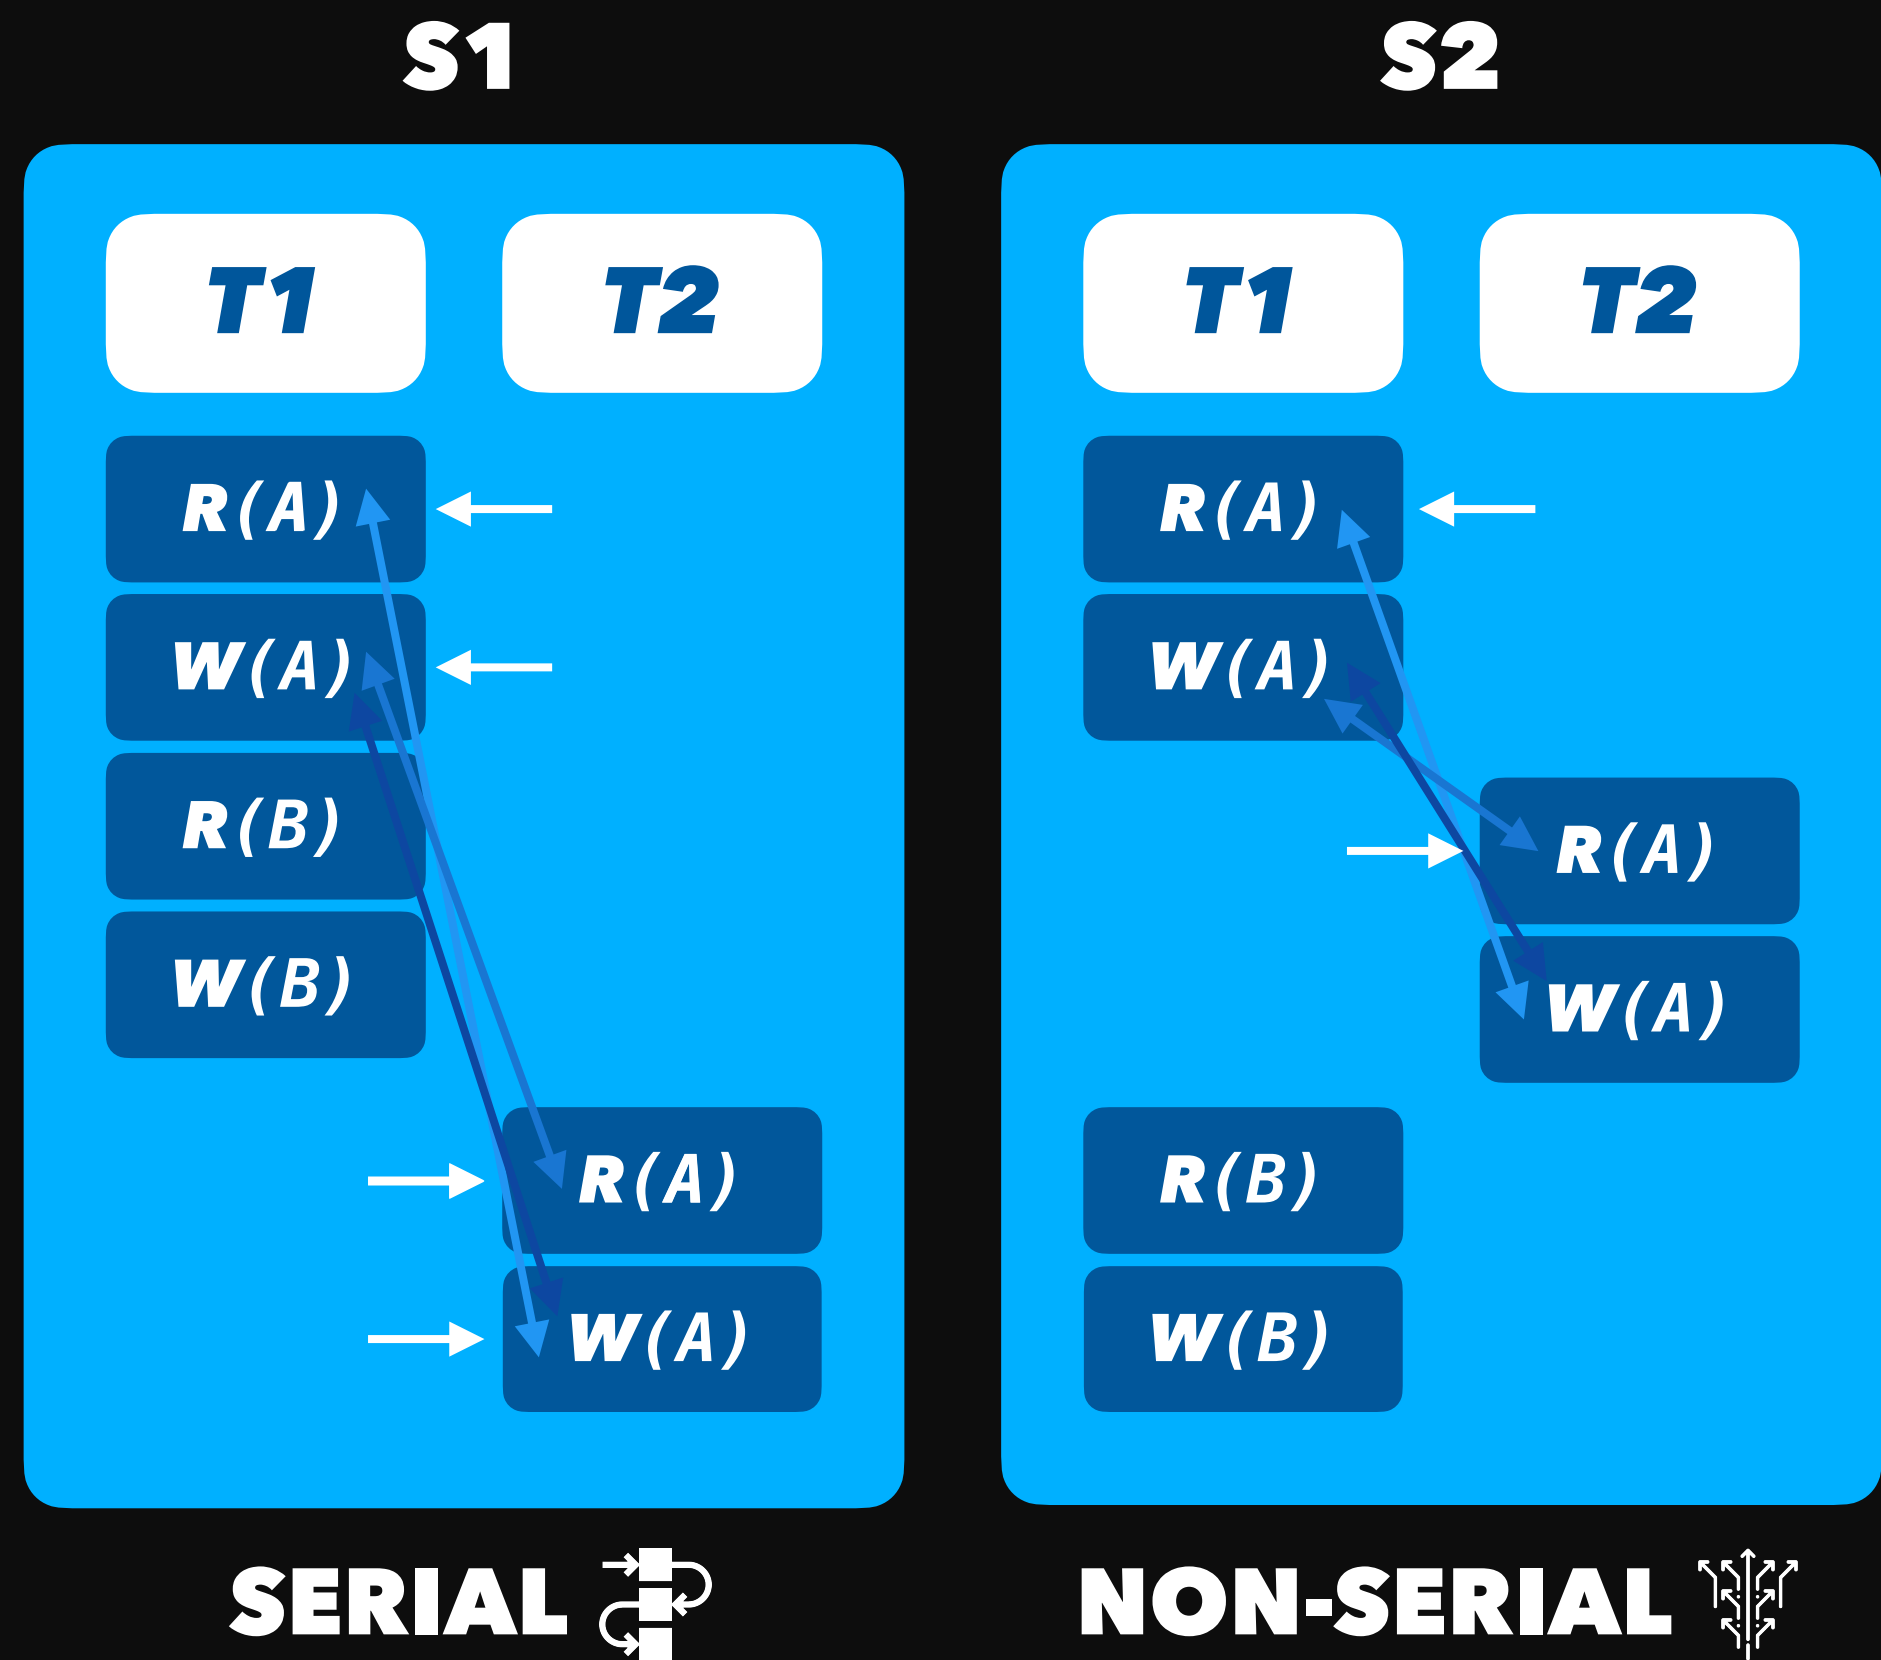
X=Y\*2

W(Y)

X: 4

Y: 12

# CONFLICT SERIALIZABLE




**Conflict operations** are on the **same data item** i.e **A, B**

**RW, WR, WW** conflicts MUST be on **different** transactions

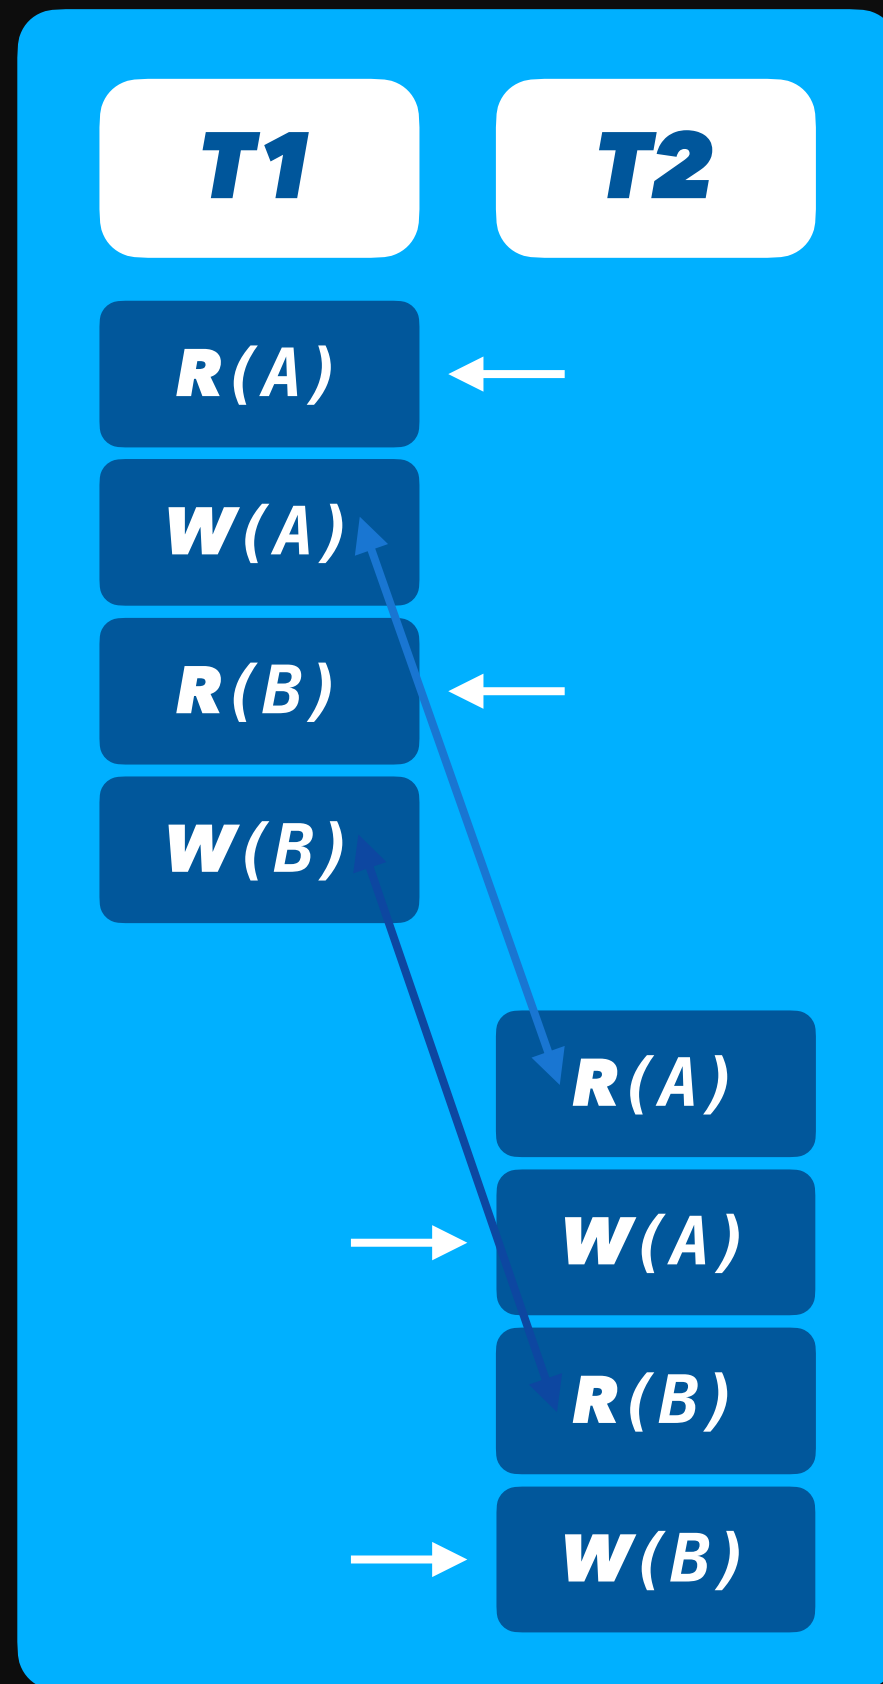
At least **1** of conflict operations MUST be a **Write**

**2 READ** operations will **not** create a **conflict**



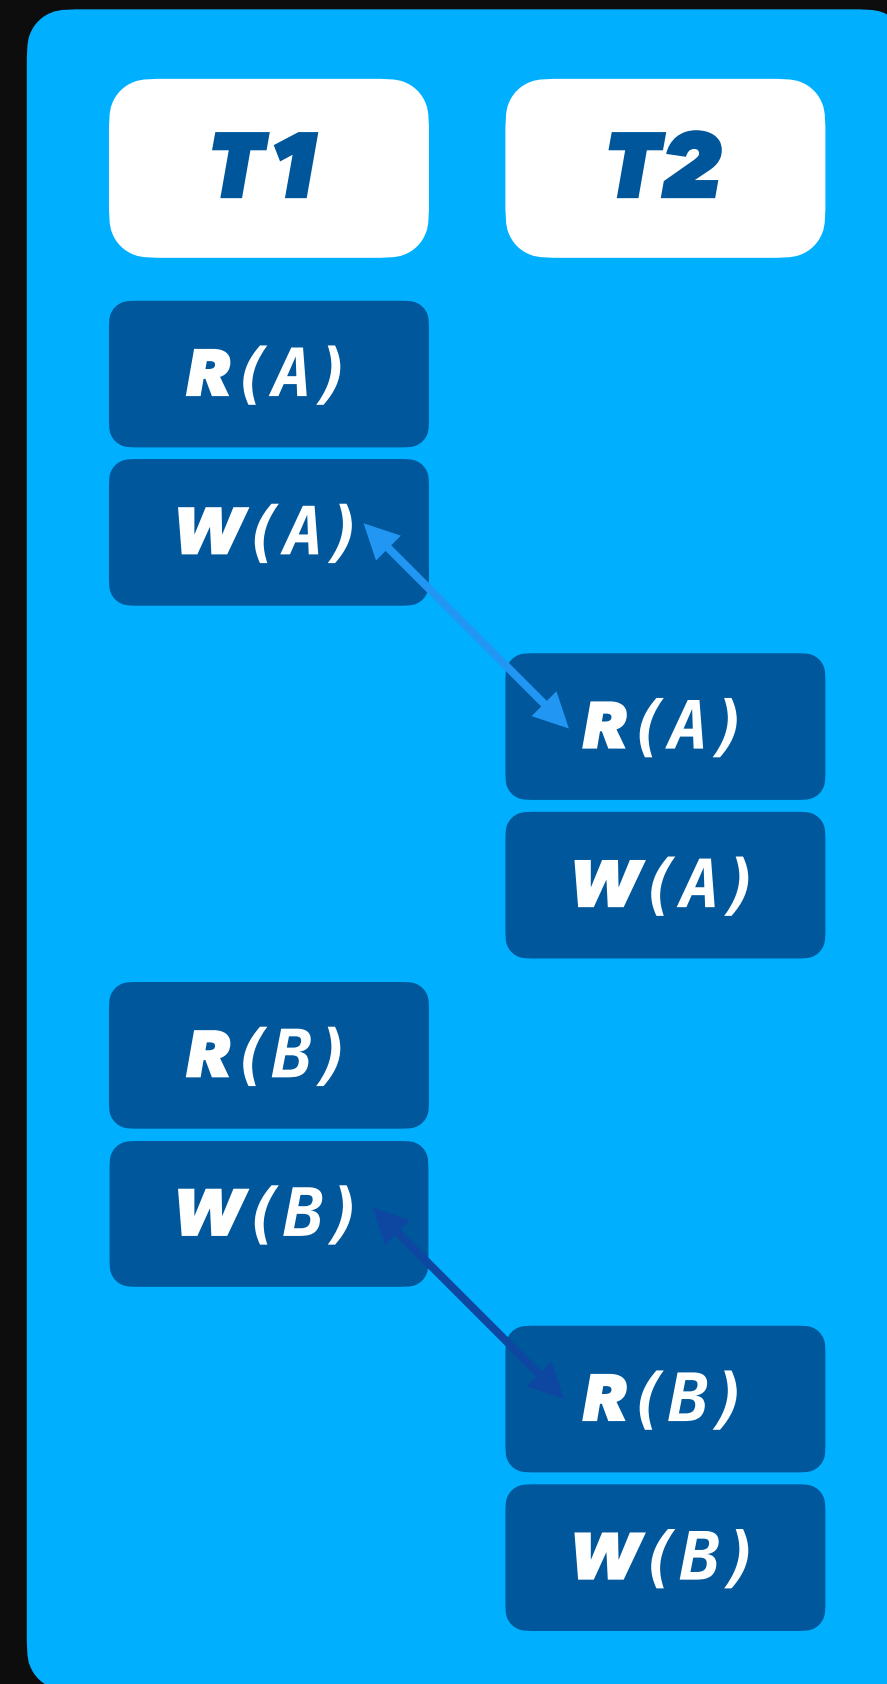
# VIEW SERIALIZABLE

S1



SERIAL 

S2



NON-SERIAL 


 *FIRST READ performed by the same transaction*  
**Initial** operations on each **data item** are on the same transaction

**Final** operations on each **data item** are on the same transaction  
 *LAST WRITE performed by the same transaction*

Number of **WR** operations MUST be the **same** on all schedules  
 *Maintain Producer-Consumer/W→R sequence*

S1

- Initial operation of data item A: T1
- Initial operation of data item B: T1
- Final operation of data item A: T2
- Final operation of data item B: T2
- Number of RW operations: 2



S2

- Initial operation of data item A: T1
- Initial operation of data item B: T1
- Final operation of data item A: T2
- Final operation of data item B: T2
- Number of RW operations: 2

# SERIALIZABILITY

LOCKING

SERIALIZATION – GRAPH CHECKING

TIMESTAMP ORDERING

COMMITMENT ORDERING

MULTIVERSION CONCURRENCY CONROL

INDEX CONCURRENCY CONROL

PRIVATE WORKSPACE MODEL

# LOCKING PROTOCOLS

**SIMPLE LOCKING**



**BASIC 2PL**

**CONSERVATIVE 2PL**

**STRICT 2PL**

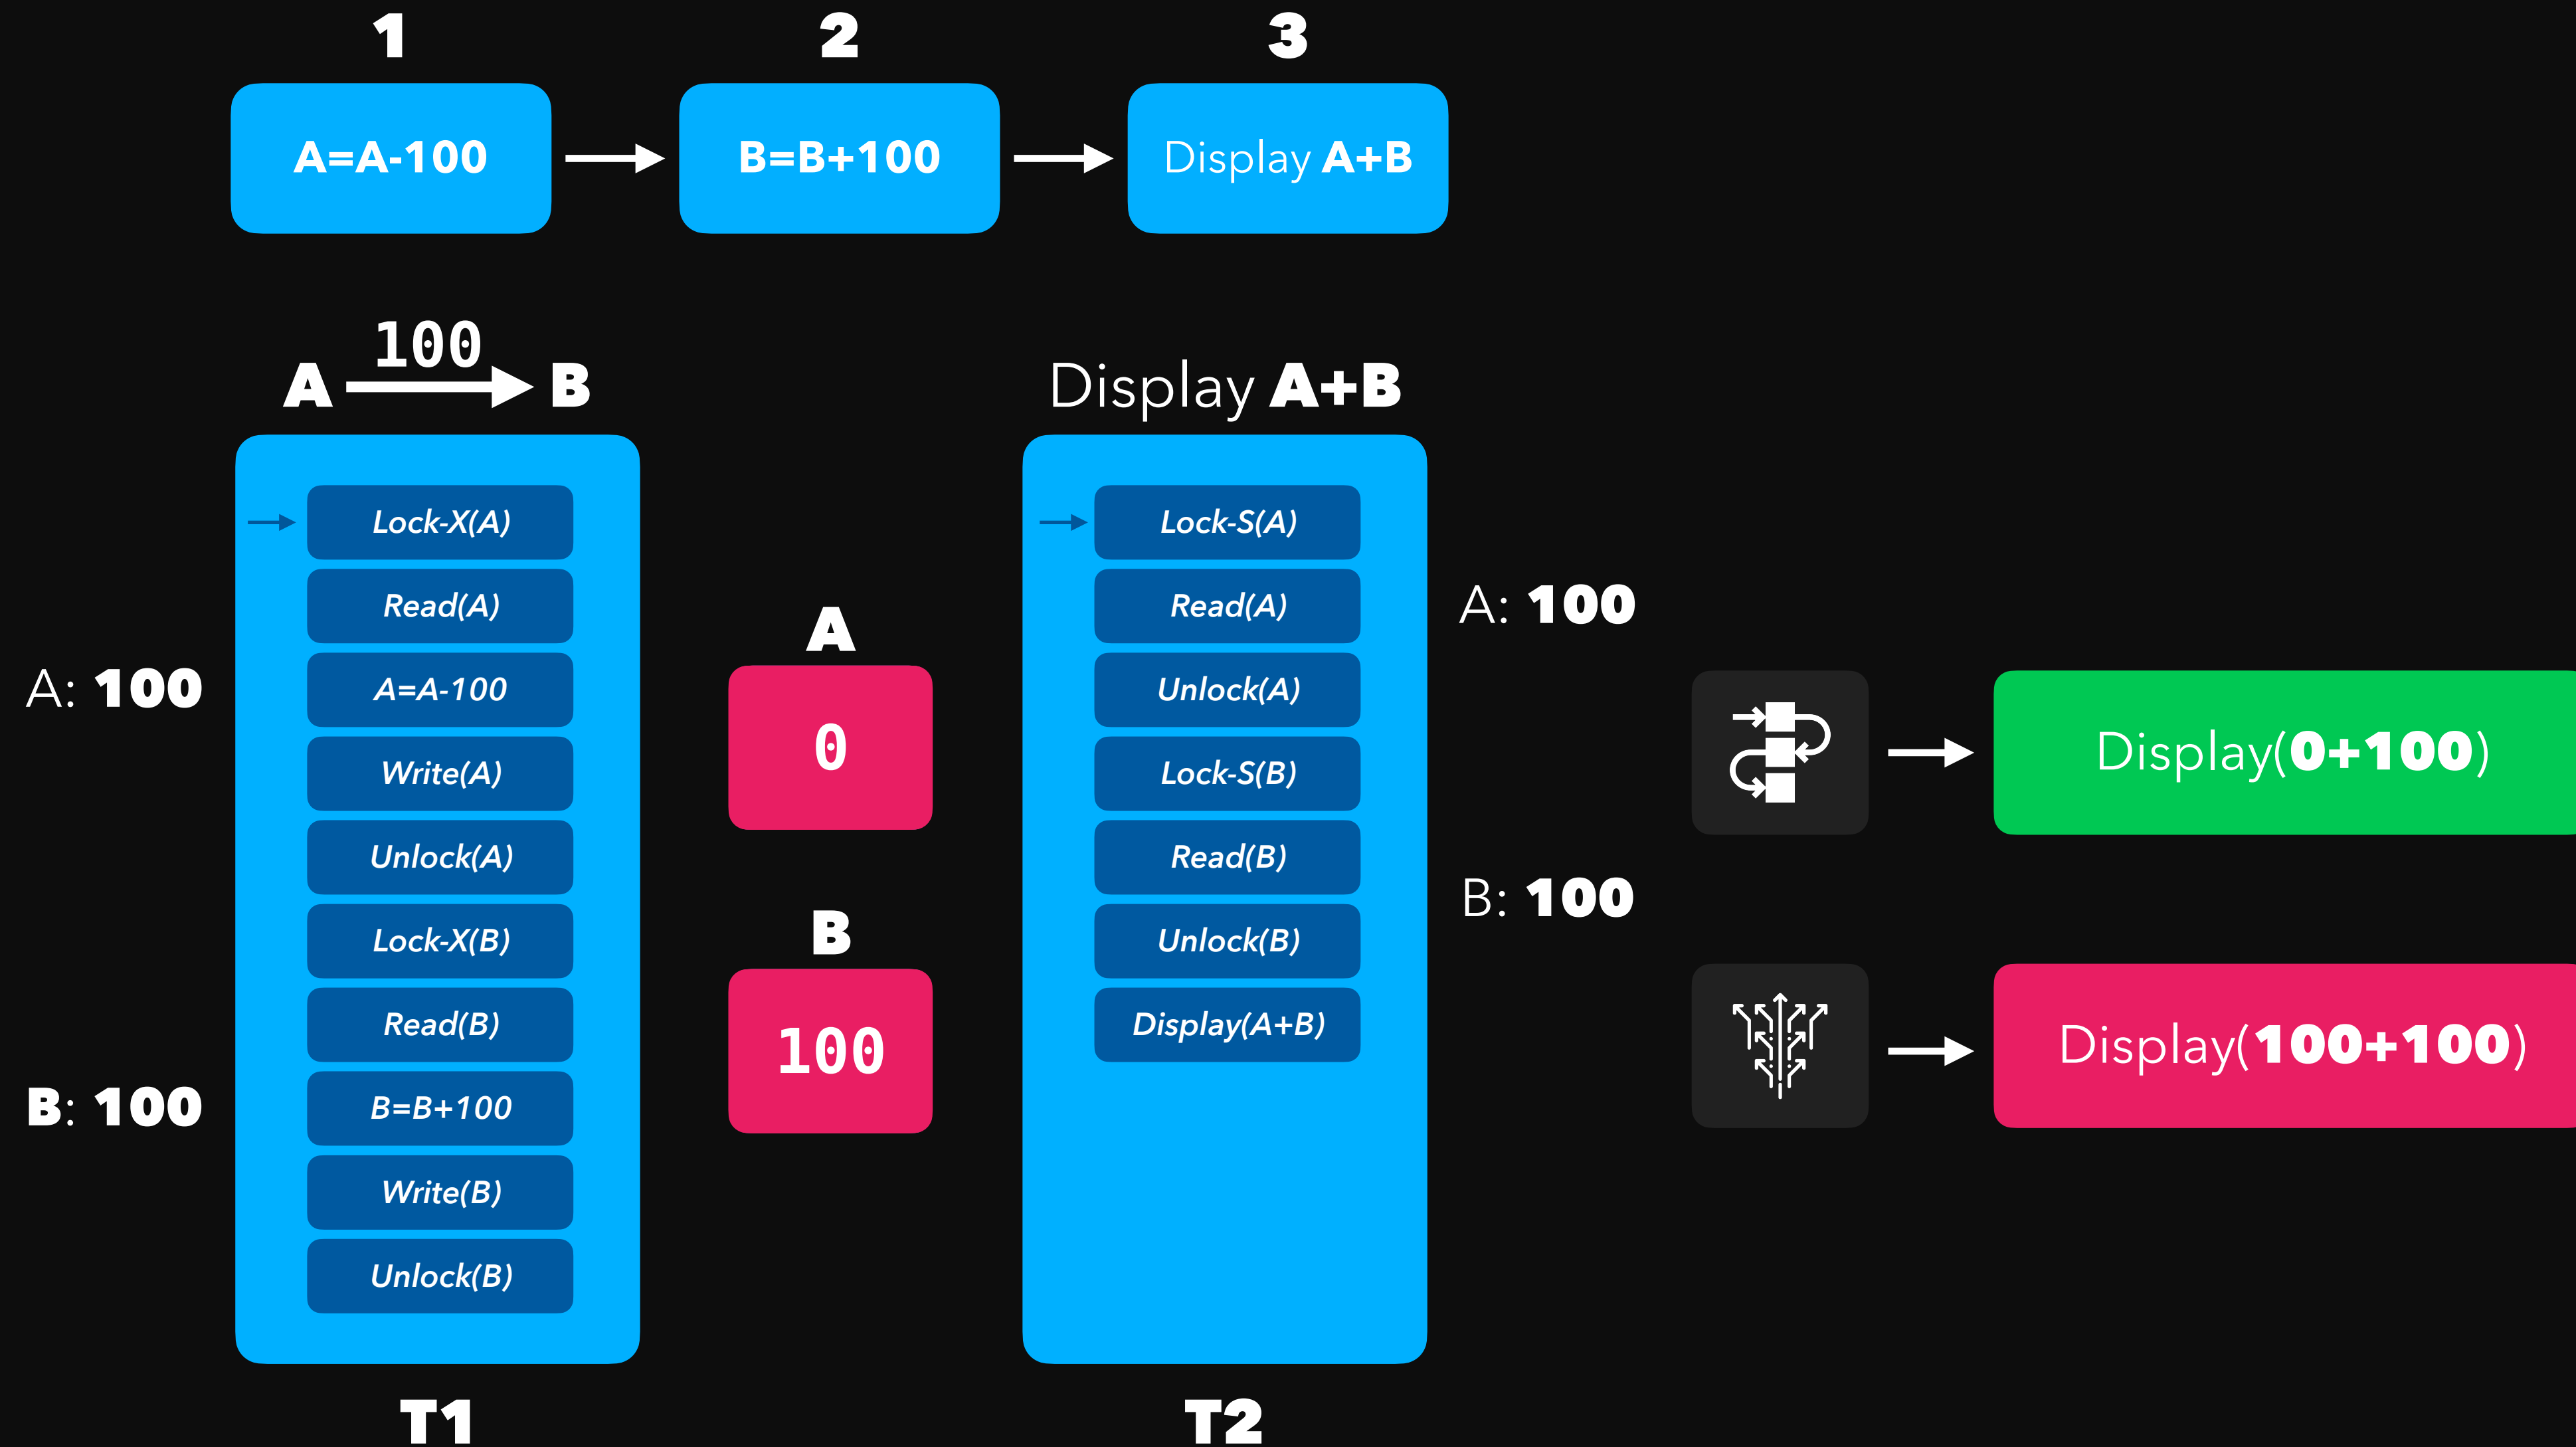
**RIGOROUS 2PL**

# SIMPLE LOCKING


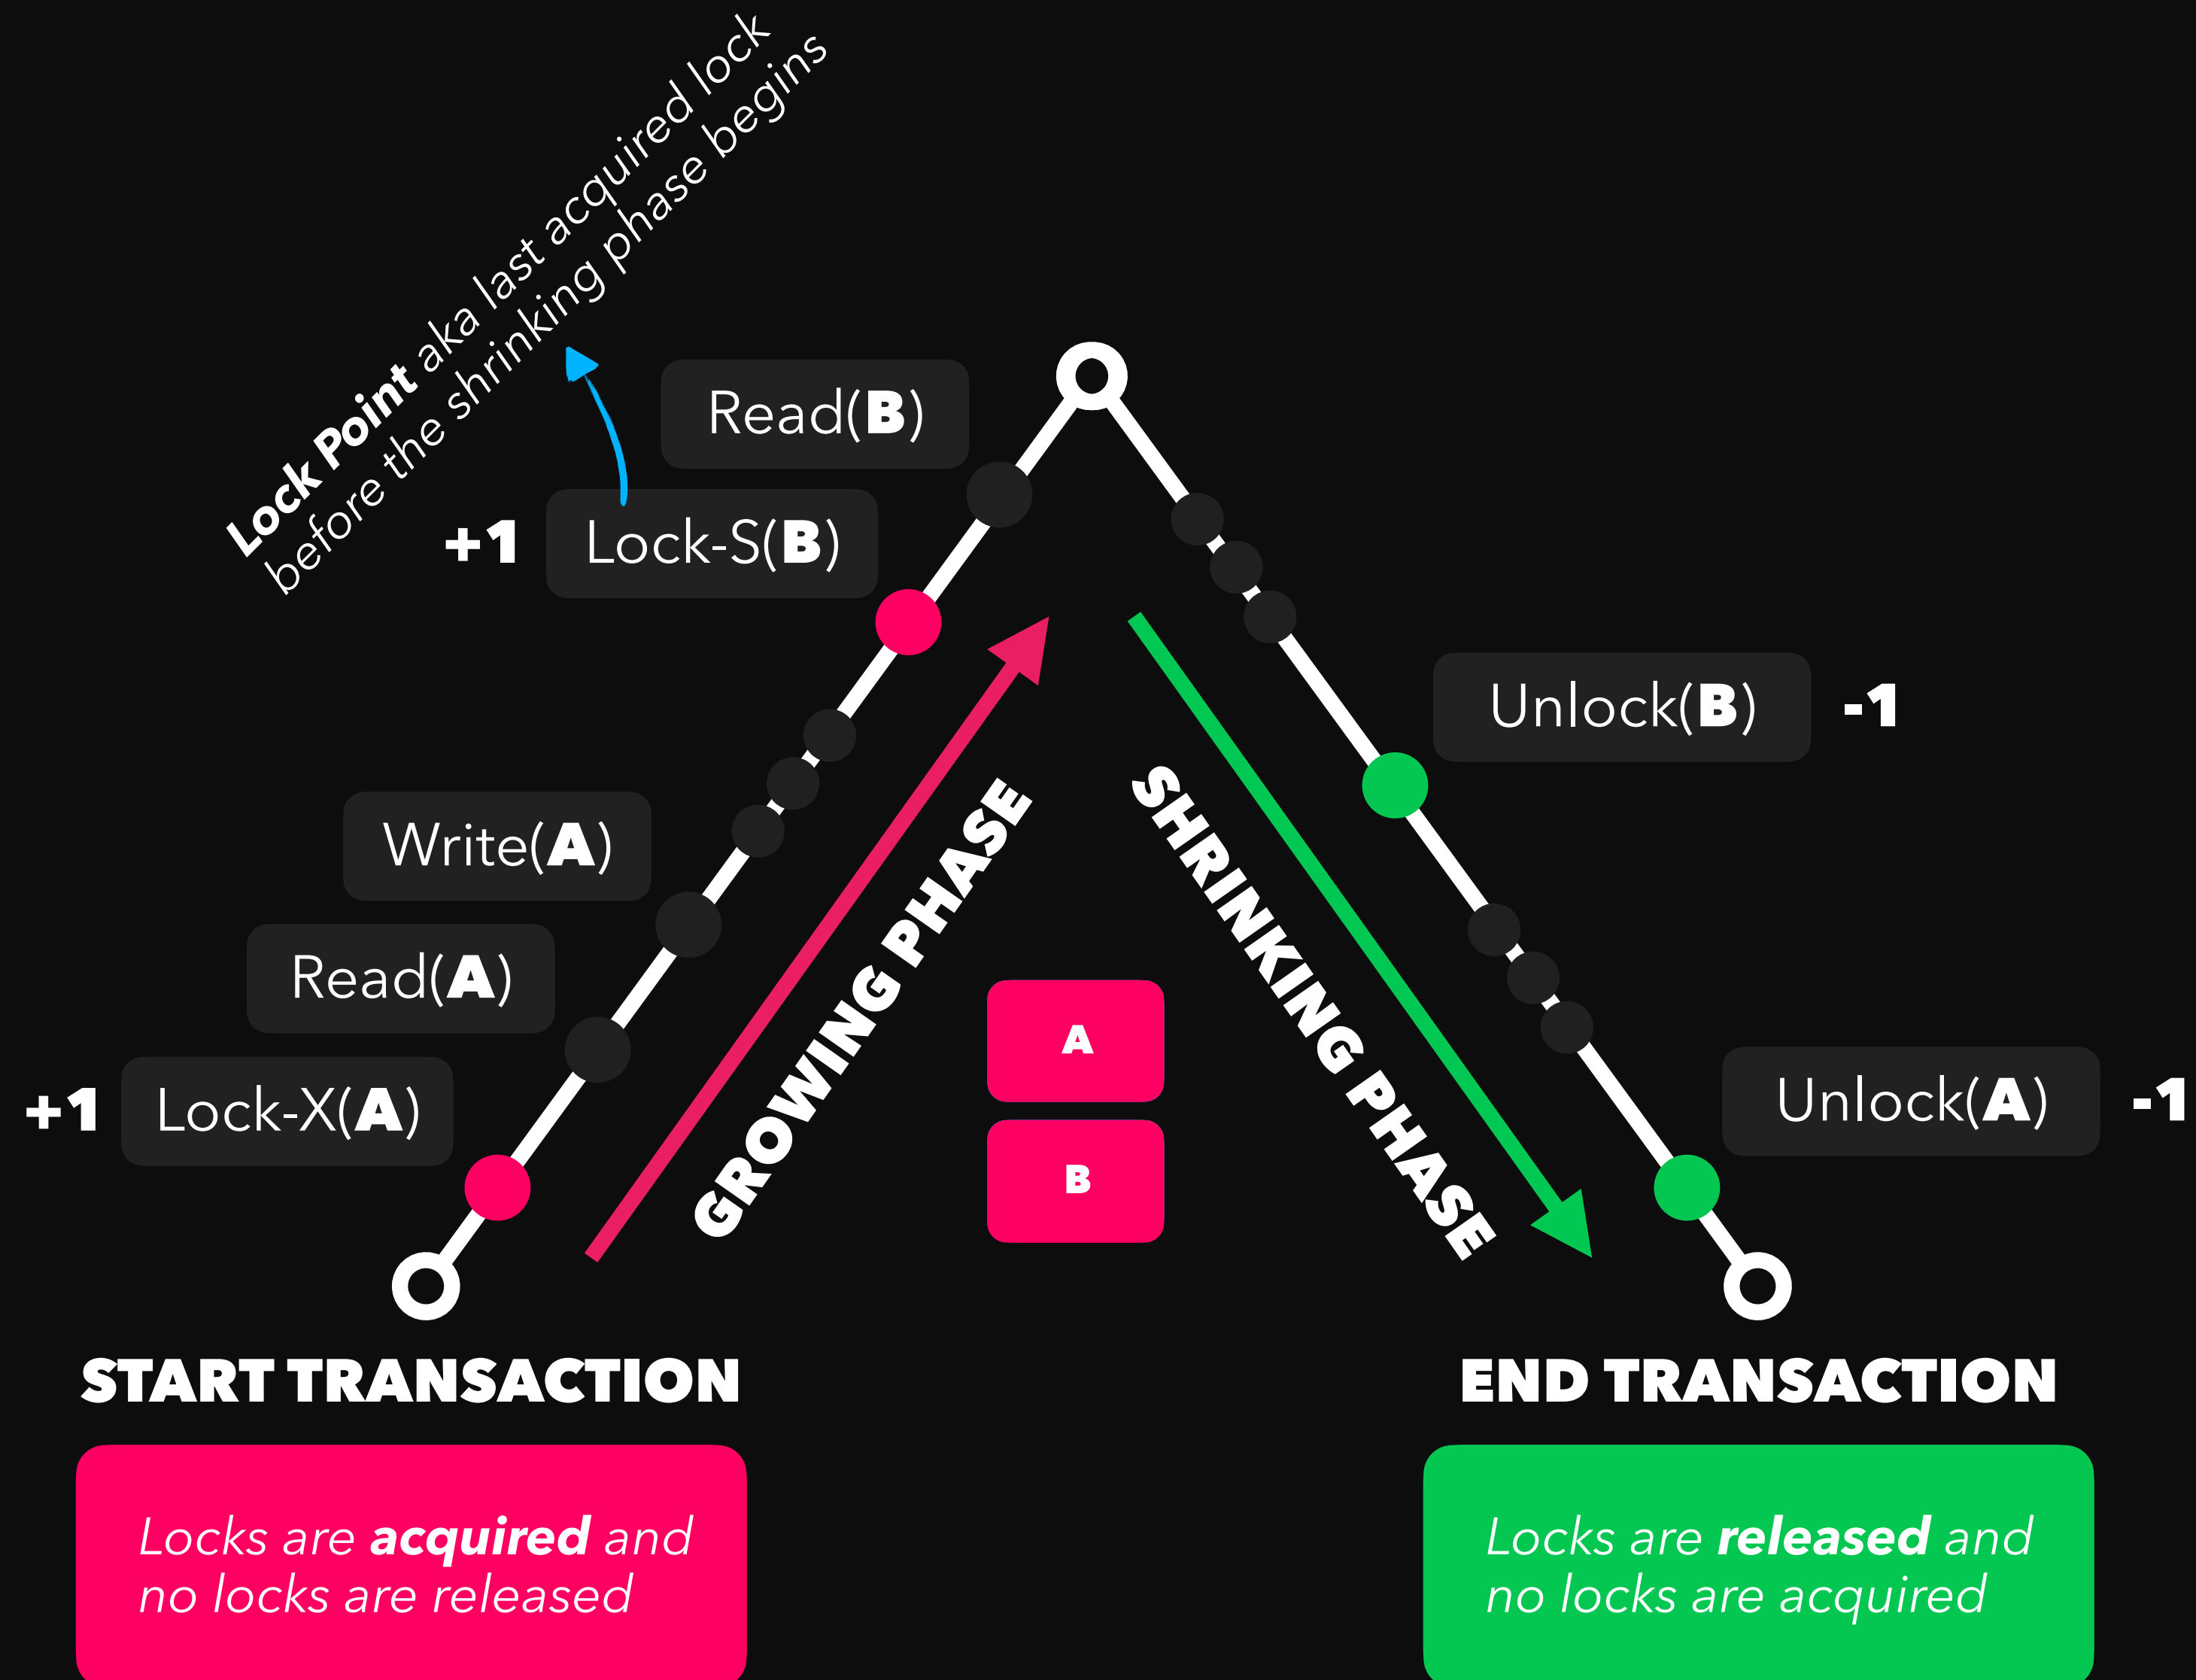


TYPE	READ LOCK	WRITE LOCK
READ LOCK	✓	X
WRITE LOCK	X	X

# SIMPLE LOCKING EXAMPLE



# BASIC 2PL




UNNECESSARY WAIT

DEADLOCKS

CASCADING ROLLBACKS



# C2PL PROTOCOL



**NO DEADLOCKS**


**DIFICULT IMPLEMENTATION**

**CASCADING ROLLBACKS**

*All locks are **acquired** before transaction start*

*Locks are **released** and no locks are acquired*

# S2PL PROTOCOL



## START TRANSACTION

Locks are **acquired** and no locks are released

## END TRANSACTION

Only **shared locks** are released

**MOST POPULAR**

**STRICT SCHEDULING**

**EASY RECOVERY**

**NO CASCADING ROLLBACKS**


**DEADLOCKS**

**COMMIT**

**ABORT**


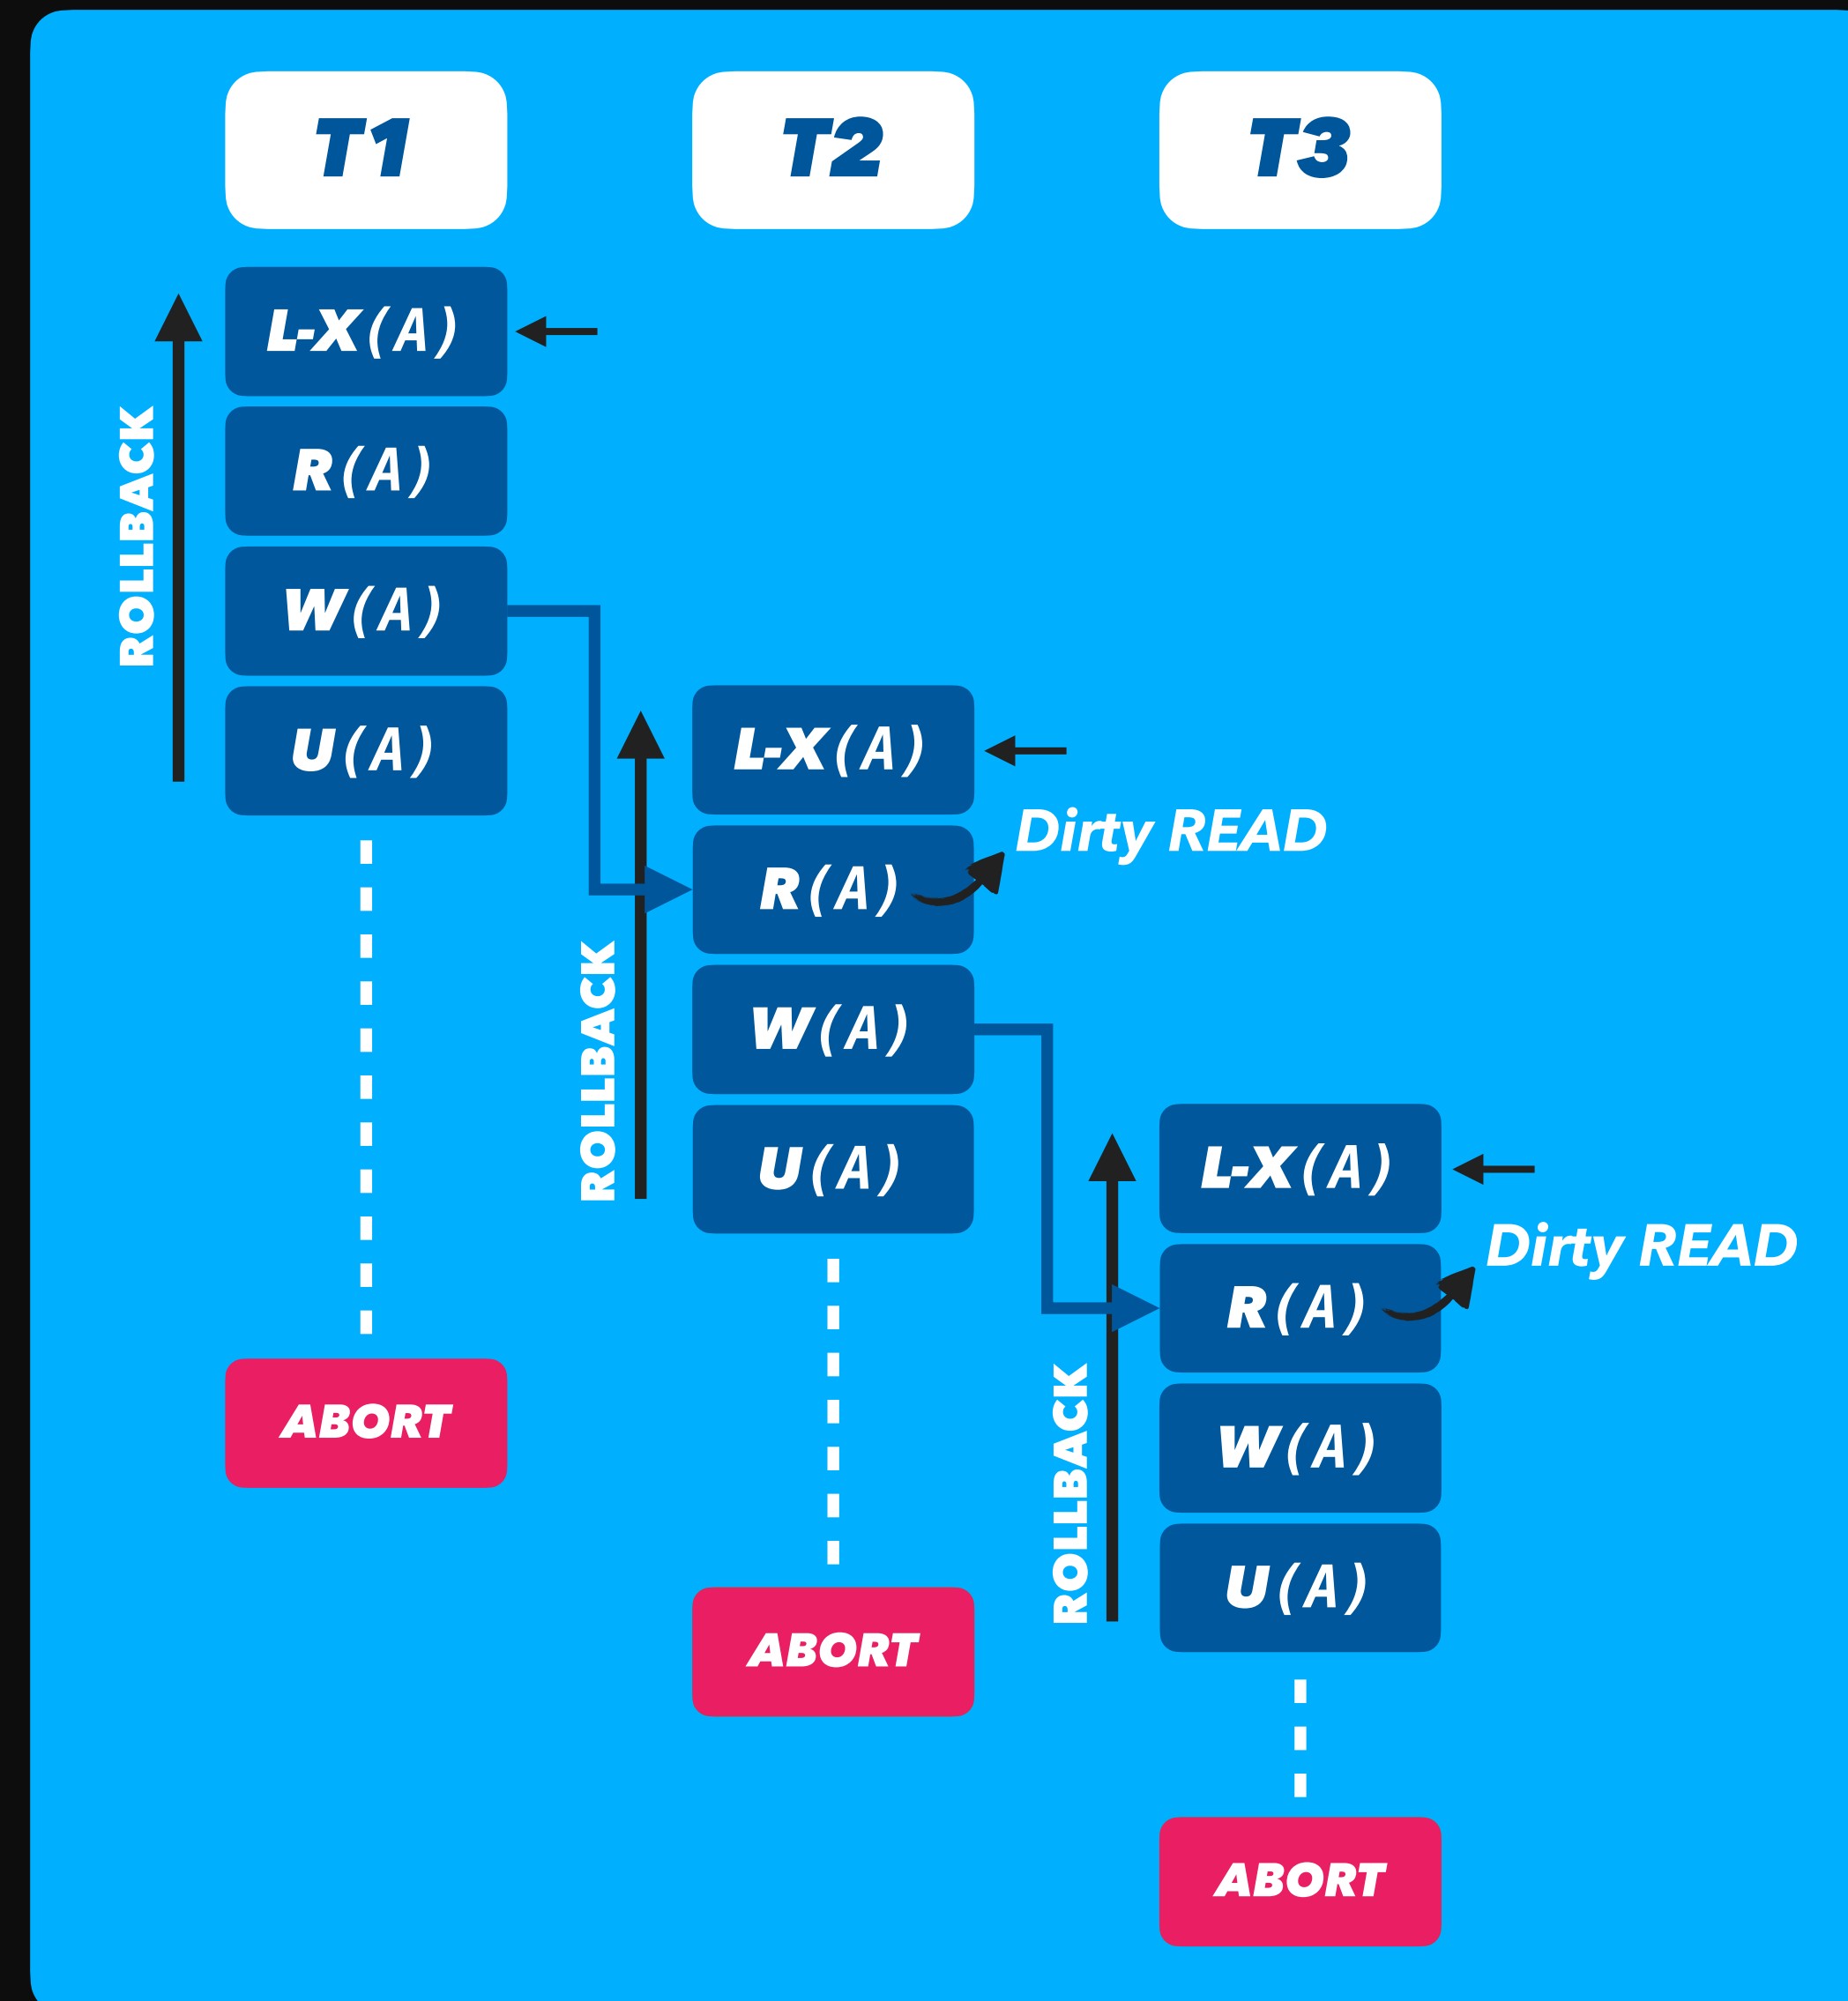
Unlock(A) -1

# SS2PL/R2PL PROTOCOL



# RECOVERABILITY

  
**CASCADING  
ABORT**



# TRANSACTION ISSUES

LOST UPDATE

DIRTY READ


INCORRECT SUMMARY

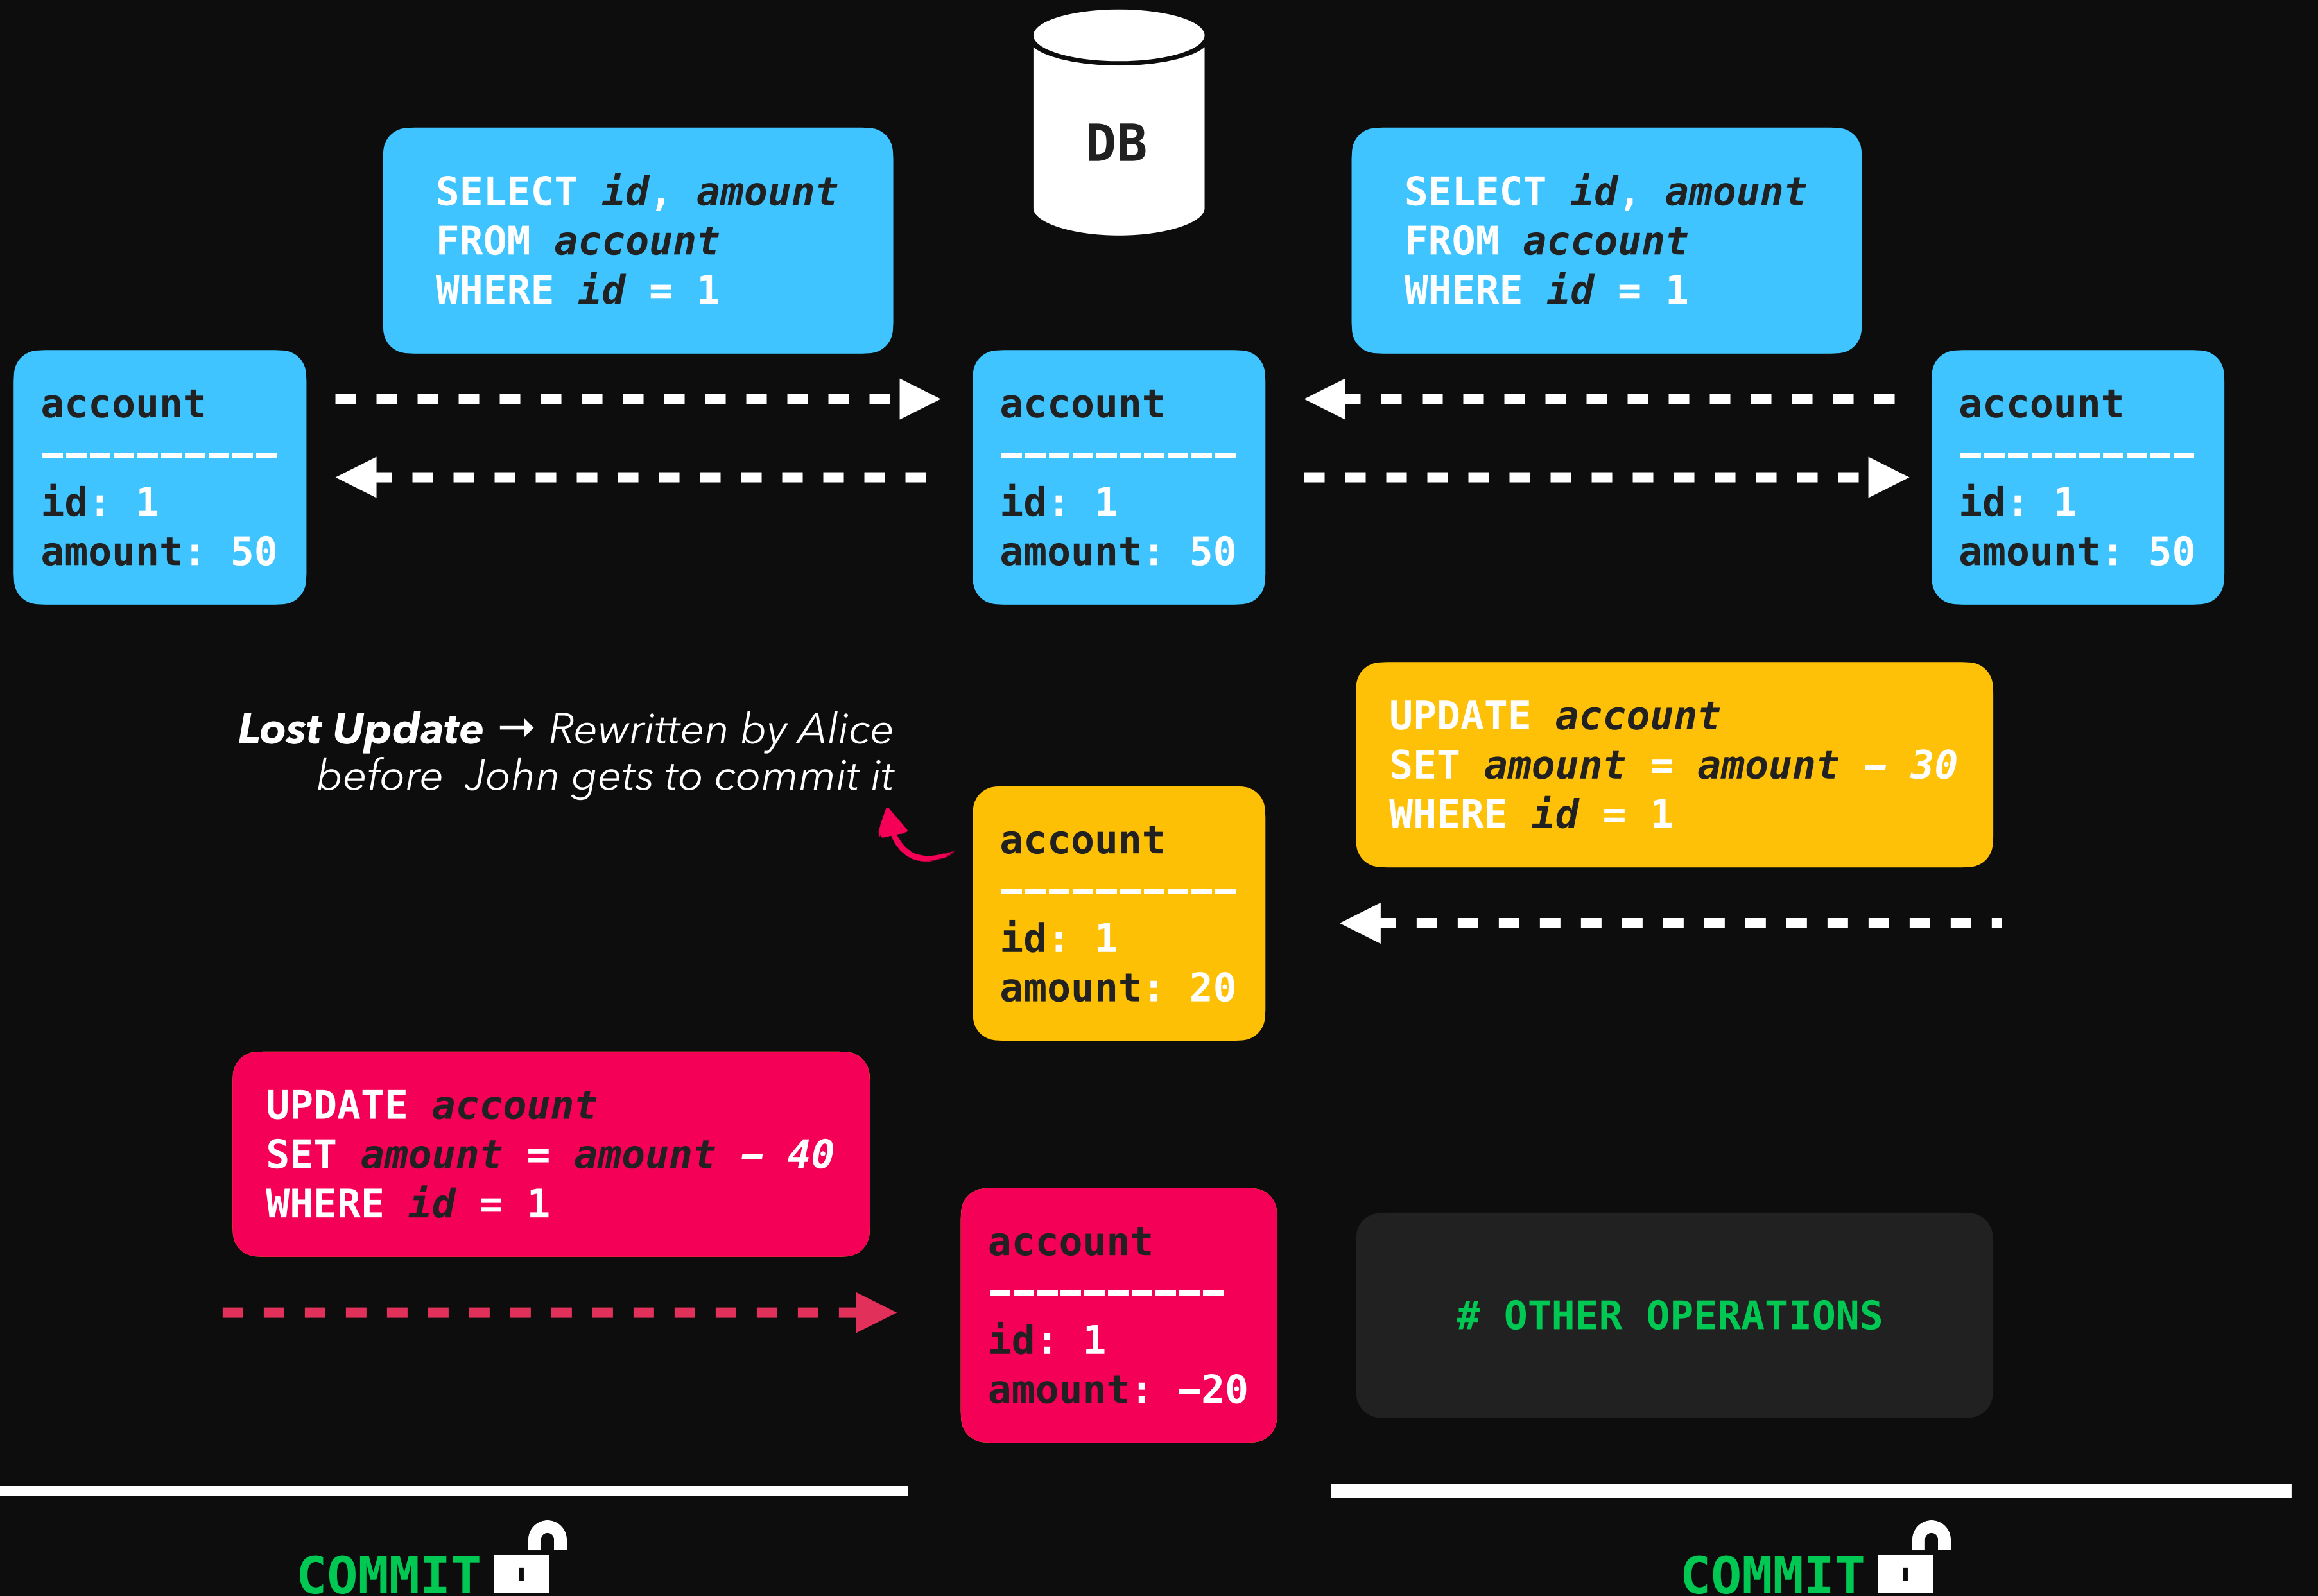
# LOCKING TYPES

OPTIMISTIC

PESSIMISTIC

SEMI-OPTIMISTIC

T1  
↑  
  
Alice





Alice

account  
-----  
id: 1  
amount: 50

# OTHER OPERATIONS

COMMIT 

SELECT *id, amount*  
FROM *account*  
WHERE *id* = 1

Shared Lock



Exclusive Lock

SELECT *id, amount*  
FROM *account*  
WHERE *id* = 1

account  
-----  
id: 1  
amount: 50

account  
-----  
id: 1  
amount: 20

UPDATE *account*  
SET *amount* = *amount* - 30  
WHERE *id* = 1

COMMIT 

account  
-----  
id: 1  
amount: 50



John






Alice

account  
-----  
id: 1  
amount: 50  
version: 1

```
UPDATE account
SET amount = amount - 40
AND version = 2
WHERE id = 1
AND version = 1
```

```
SELECT id, amount
FROM account
WHERE id = 1
```



account  
-----  
id: 1  
amount: 50  
version: 1

account  
-----  
id: 1  
amount: 30  
version: 2

```
SELECT id, amount
FROM account
WHERE id = 1
```

account  
-----  
id: 1  
amount: 50  
version: 1

```
UPDATE account
SET amount = amount - 30
AND version = 2
WHERE id = 1
AND version = 1
```



John

NO UPDATE  
ROW MATCHED

ABORT 

COMMIT 